

# ZFS on FreeBSD: A Quick Tutorial

Tai-hwa Liang  
<avatar@FreeBSD.org>

# History/Availability

---

- Designed by Sun Microsystems (Sep-14-2004).
- First OpenSolaris inclusion: Nov-16-2005.
- Solaris 10 inclusion: Jun-2006.
- Became part of FreeBSD on Apr-06-2007.
  - Ported by Pawel Jakub Dawidek <pjd@FreeBSD.org>.
- Mac OS X 10.5 Leopard (?)
- Linux (via FUSE)
  - CDDL is not 'compatible' with GPL

# ZFS Features

---

## ☐ Pooled Storage

- Build-in volume manager (LVM-like) and RAID

## ☐ Transactional Semantics

- Think about databases...
- No more fsck, journaling

## ☐ Lightweight filesystem creation

- Snapshot (rollback to "undo" anything bad)
- Clones (writable snapshots)

## ☐ Mind-boggling Scalability

- 128 bits(256 quadrillion zettabytes)
  - ▷ giga -  $10^9$
  - ▷ zetta -  $10^{21}$

# ZFS Features (cont.)

---

- ❑ Data Integrity
  - Checksumming
  - Self healing
  
- ❑ Platform Independent Endianness
  
- ❑ Built-in Compression
  - lzjb
  - gzip-N
  
- ❑ Simplified Administration
  - No more LVM
  - Delegated to non-root users

# ZFS Limitations

---

- ❑ Per-user or per-group quota is not supported.
  - Workaround: create user-owned filesystem and set its own size limit with 'zfs set quota=XXG'
  
- ❑ Transparent encryption is not supported, yet.
  - WIP: <http://www.opensolaris.org/os/project/zfs-crypto>
  - Workaround: geli(8)
  
- ❑ It's not possible to change the number of disk in a RAID-Z set.
  
- ❑ It's not possible to reduce the number of vdevs in a zpool.
  - WIP: <http://www.opensolaris.org/jive/thread.jspa?messageID=118280>

# What About The Performance?

---

- Testing environment:

- 7.0-CURRENT-200706(snapshot) GENERIC kernel
- Intel Core 2 T5500 @ 1.66GHz
- 1GB memory

- Extract a 447MB tarball(-CURRENT source tree as of Aug-21-2007):

- UFS + SoftUpdate: 2:45.86
- ZFS: 1:06.00

- 'make buildworld buildkernel':

- UFS + SoftUpdate: 1:48:11.96
- ZFS: 2:21:16.75

# Host with 'Smaller' Memory

---

- ❑ 512MB memory(it's strongly discouraged to use ZFS without at least 1 GB of memory.)

- /boot/loader.conf:

- vm.kmem\_size="256M"
- vfs.zfs.prefetch\_disable="1"

- /etc/sysctl.conf:

- kern.maxvnodes=22500

- ❑ 1024MB memory

- /boot/loader.conf:

- vfs.zfs.arc\_max="100M"
- vm.kmem\_size\_max="600M"
- vfs.zfs.prefetch\_disable="1"

- /etc/sysctl.conf:

- kern.maxvnodes=50000

# Troubleshooting

---

- ❑ For ZFS-on-root setup can't mount ZFS root:
  - /boot/zfs/zpool.cache must be up-to-date on the boot filesystem(zfs import).
  
- ❑ Ordinary users can not create snapshot?
  - Make sure that they are in the 'operator' group.
  
- ❑ panic: kmem\_alloc(131072): kmem\_map too small: xxxxx total allocated
  - Not recommended for host < 1GB of memory
    - [http://people.freebsd.org/~pjd/patches/vm\\_kern.c.2.patch](http://people.freebsd.org/~pjd/patches/vm_kern.c.2.patch)
  - Set vm.kmem\_size and vm.kmem\_size\_max



# References

---

## □ Live Demo

- <http://people.freebsd.org/~pjd/misc/zfs>

## □ Guides

- Solaris ZFS Administration Guide
  - ▷ <http://opensolaris.org/os/community/zfs/docs/zfsadmin.pdf>
- ZFS - FreeBSD Wiki
  - ▷ <http://wiki.freebsd.org/ZFS>

## □ NFS and ZFS, a fine combination(read before you want to set zil\_disable=1)

- [http://blogs.sun.com/roch/entry/nfs\\_and\\_zfs\\_a\\_fine](http://blogs.sun.com/roch/entry/nfs_and_zfs_a_fine)

## □ FreeBSD/ZFS - last word in operating/file systems

- [http://people.freebsd.org/~pjd/pubs/eurobsdcon07\\_zfs.pdf](http://people.freebsd.org/~pjd/pubs/eurobsdcon07_zfs.pdf)