



Open Source Tools for Grid Application Development

Hurng-Chun Lee, Academia Sinica and CERN

Hurng-Chun.Lee@cern.ch



Outline

- Grid computing in a nutshell
- Grid application development
- Open source tools for grid application



It's a computing infrastructure ...

“A computational grid is a hardware and software infrastructure that provides dependable, consistent, pervasive, and inexpensive access to high-end computational capabilities.”

- C. Kesselman and I. Foster (1998)



It's about resource sharing and collaboration ...

Grid computing is concerned with “coordinated resource sharing and problem solving in dynamic, multi-institutional **virtual organization**.”

- The Anatomy of the Grid (2000)



It should be “open” and “high-level”

- Grid is a system ...
 - ... coordinating resources that are not subject to centralized control
 - ... using standard, open, general-purpose protocols and interfaces
 - ... delivering nontrivial qualities of services

- I. Forster (2002)



So, Grid is ...

- a service for sharing computer power and data storage capacity over the Internet
 - Grid Café (2003)



So, Grid is ...

- a service for sharing computer power and data storage capacity over the Internet
 - Grid Café (2003)

World-Wide Web ⇒ Share of information

World-Wide Grid ⇒ Share of computing power



The dream - a global computer

- a system goes well beyond simple communication between computers, and aims ultimately to turn the global network of computers into one vast computational resource

The dream - a global computer

- a system goes well beyond simple communication between computers, and aims ultimately to turn the global network of computers into one vast computational resource

Here is your input data coming from



Here is your output data going to

Here is your process is run

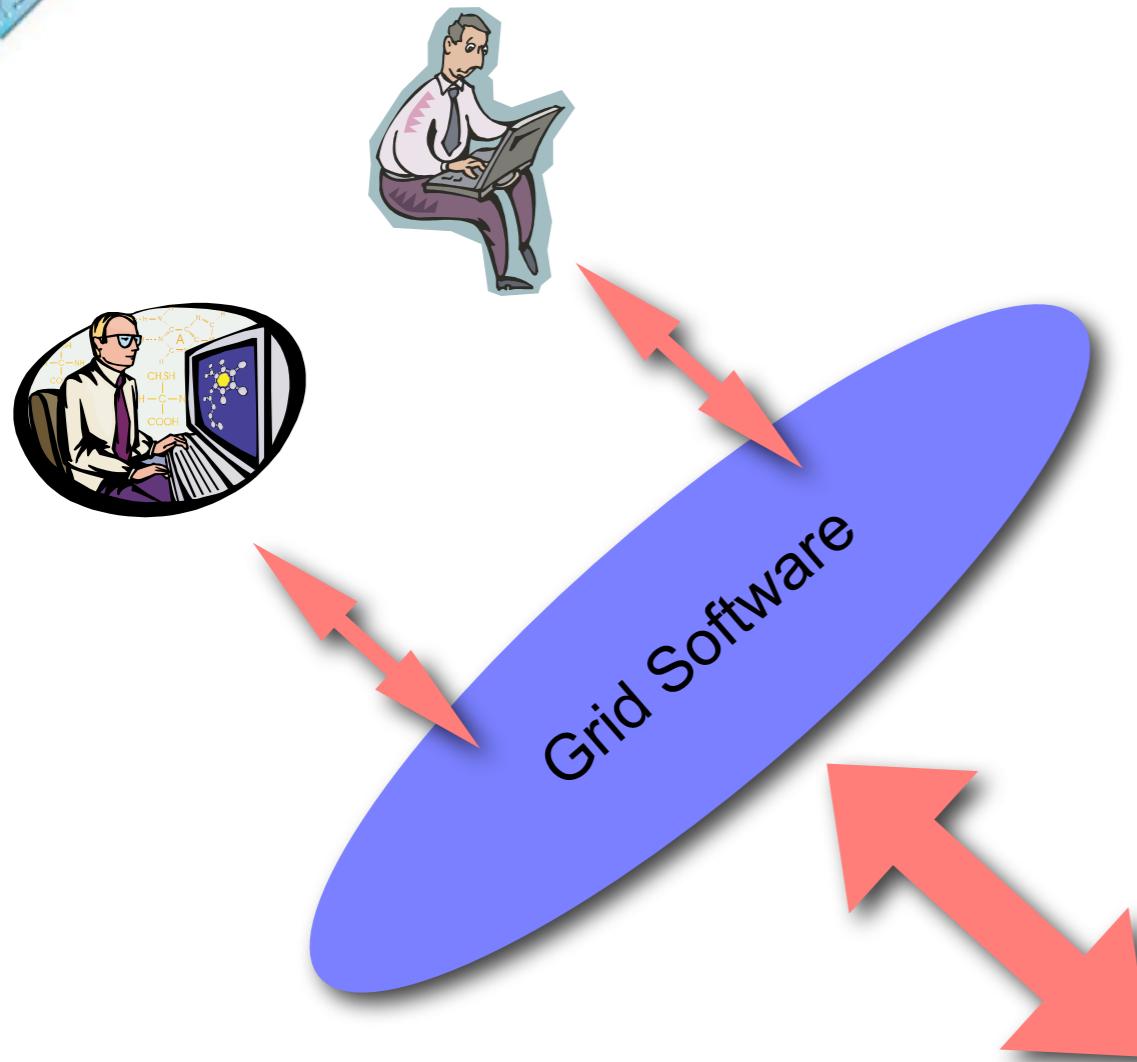
You are here



The reality ...

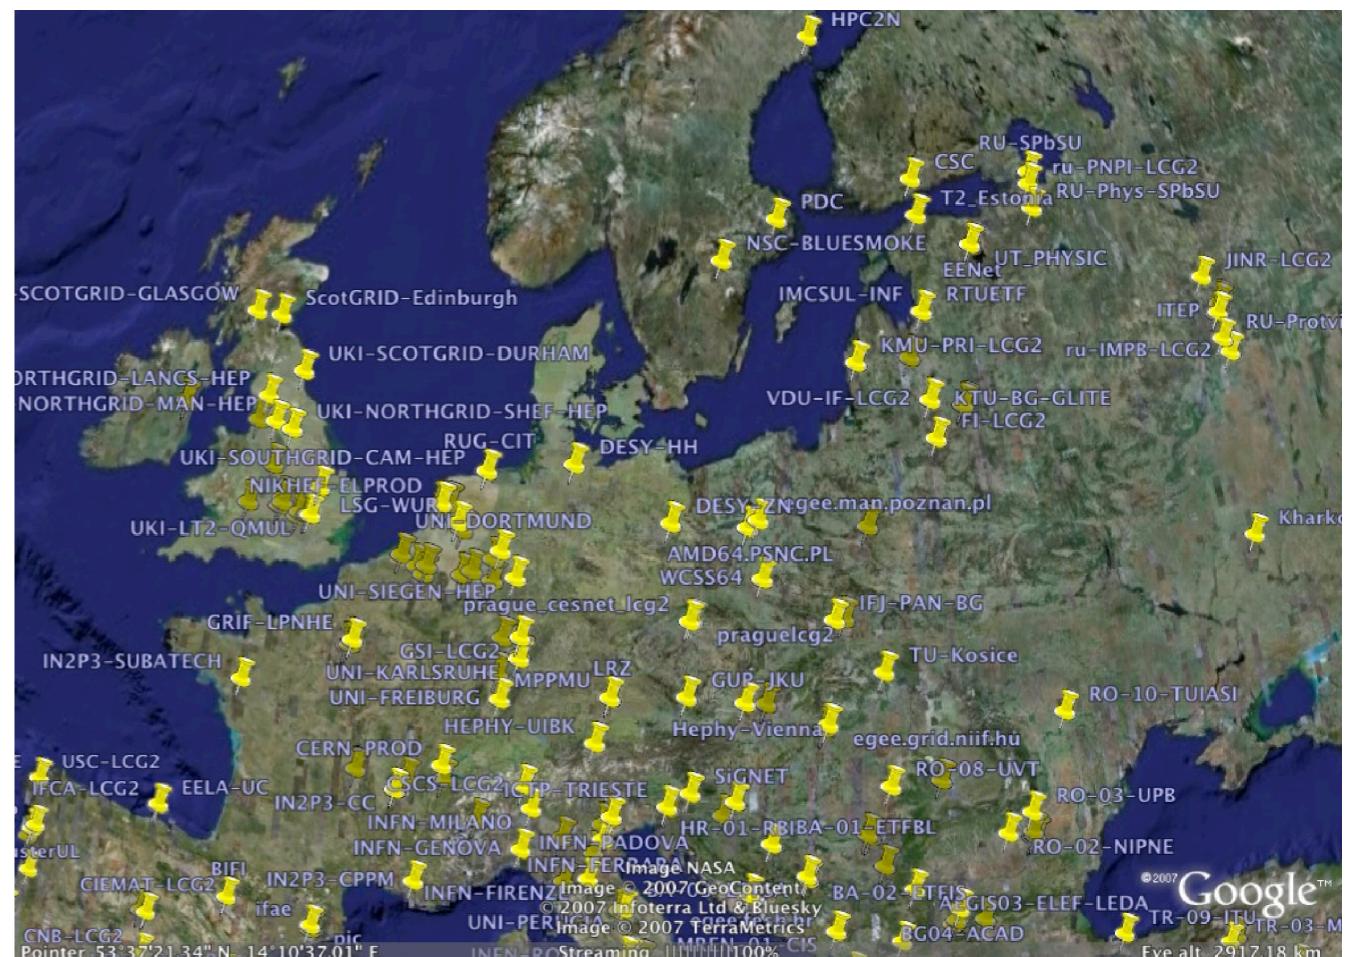


Building the Grid = software development



- Driven by scientific challenges
 - High-energy physics
 - Biomedical applications
 - etc.

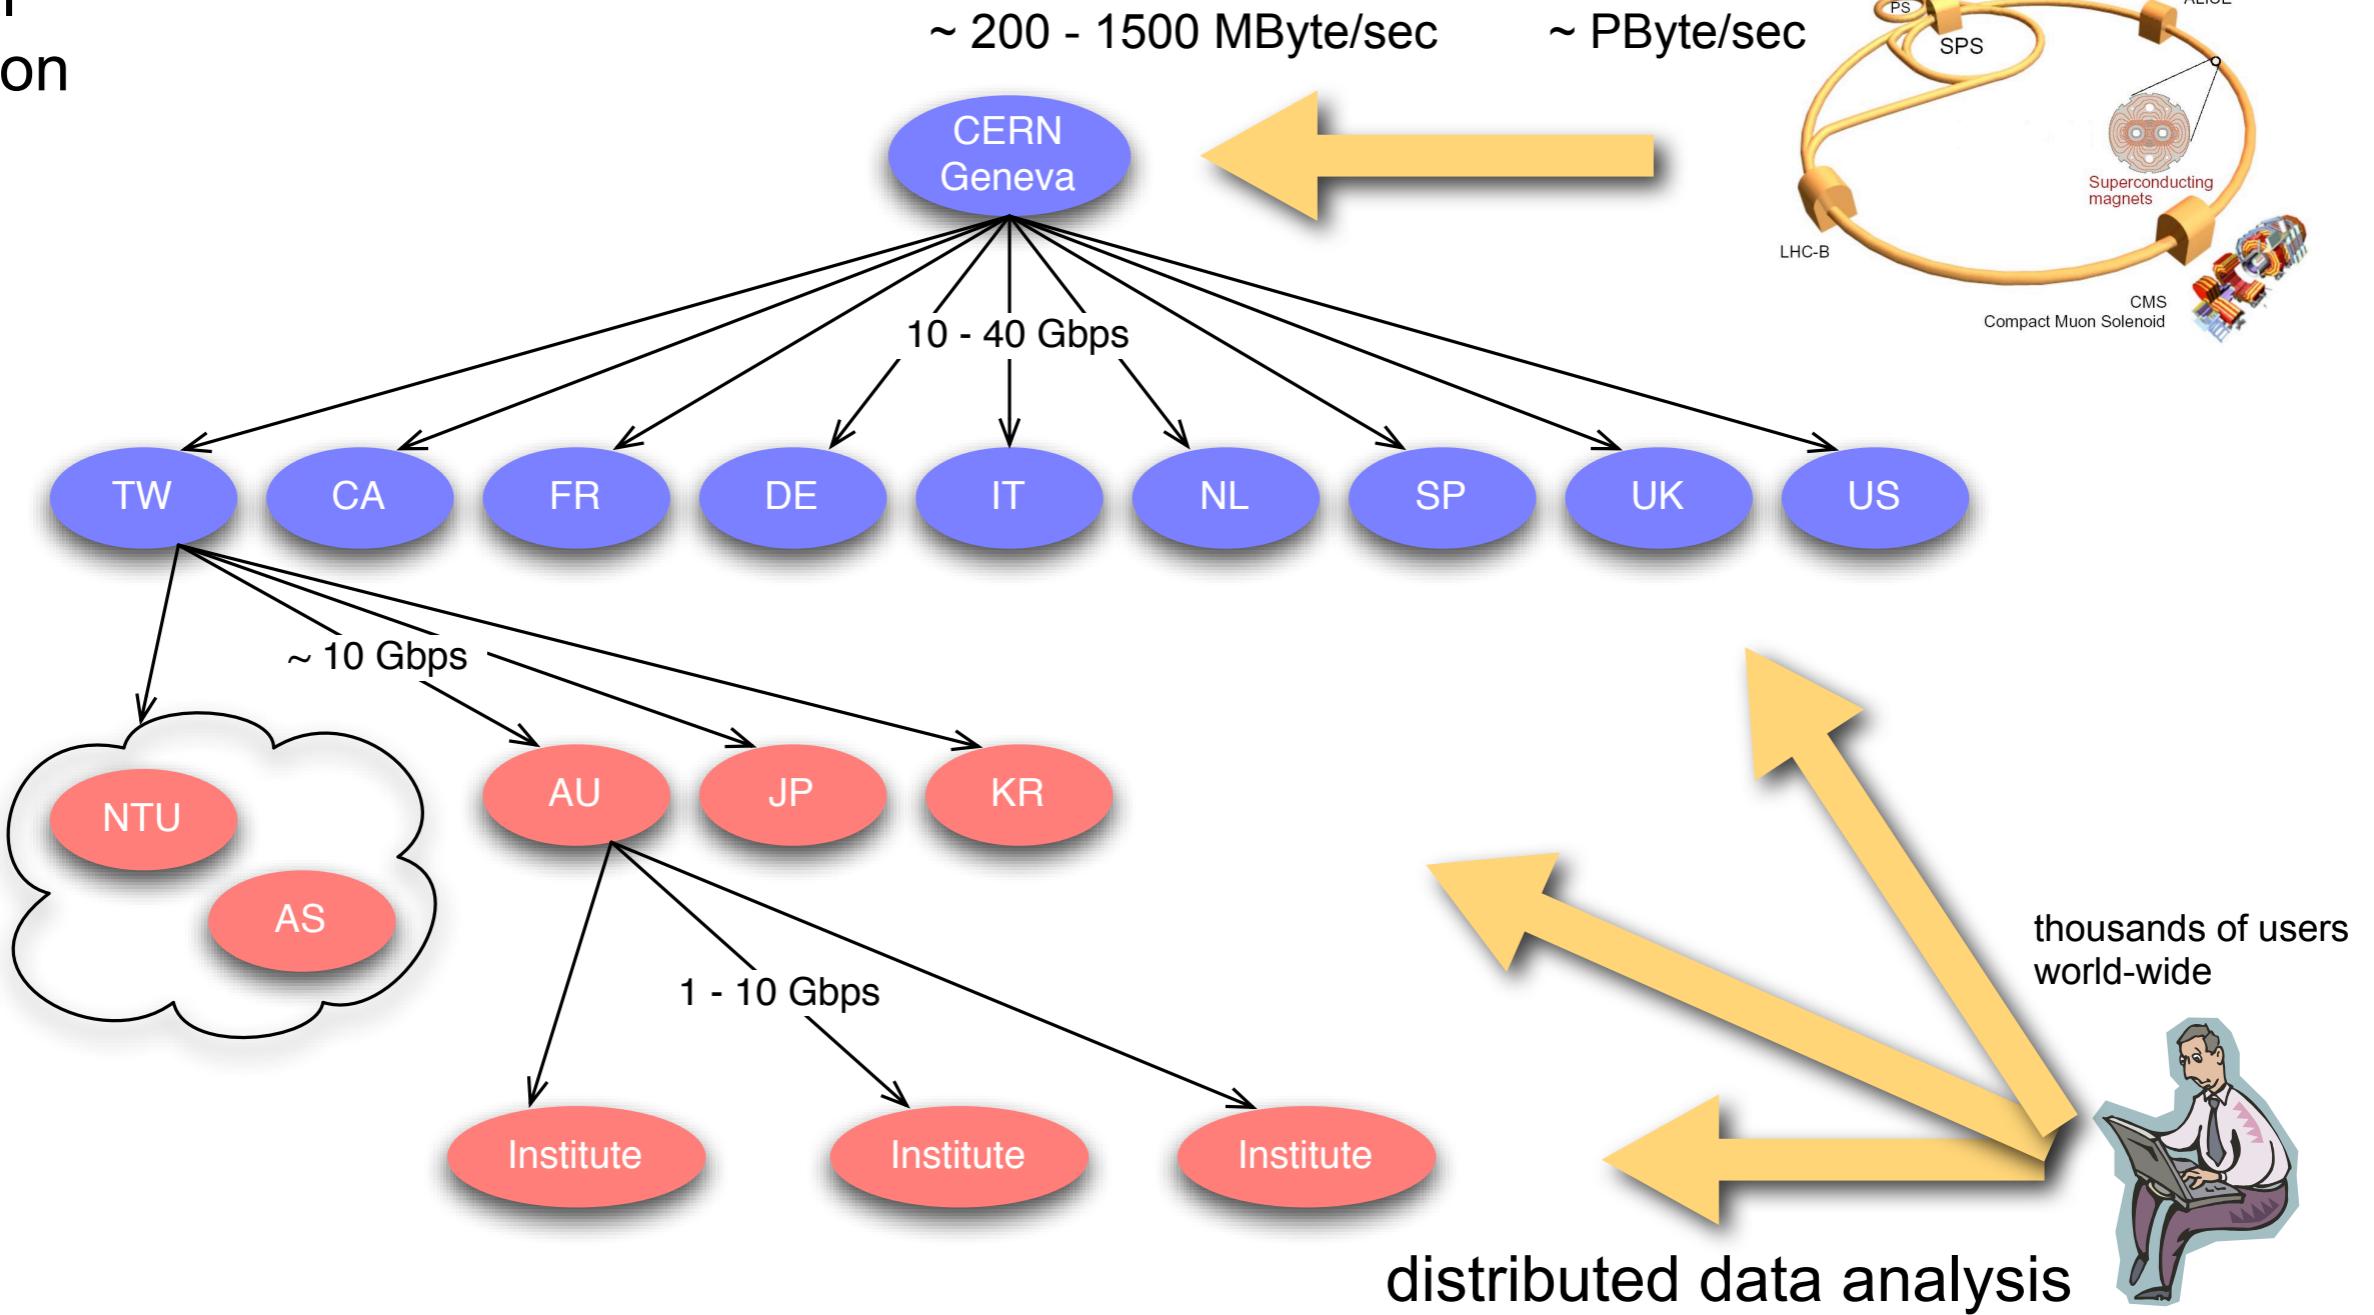
- build on top of the existing computing infrastructure (fabric)
 - internet
 - computing/data centers
 - etc.





Use case: High-energy physics

central production





Use case: High-energy physics

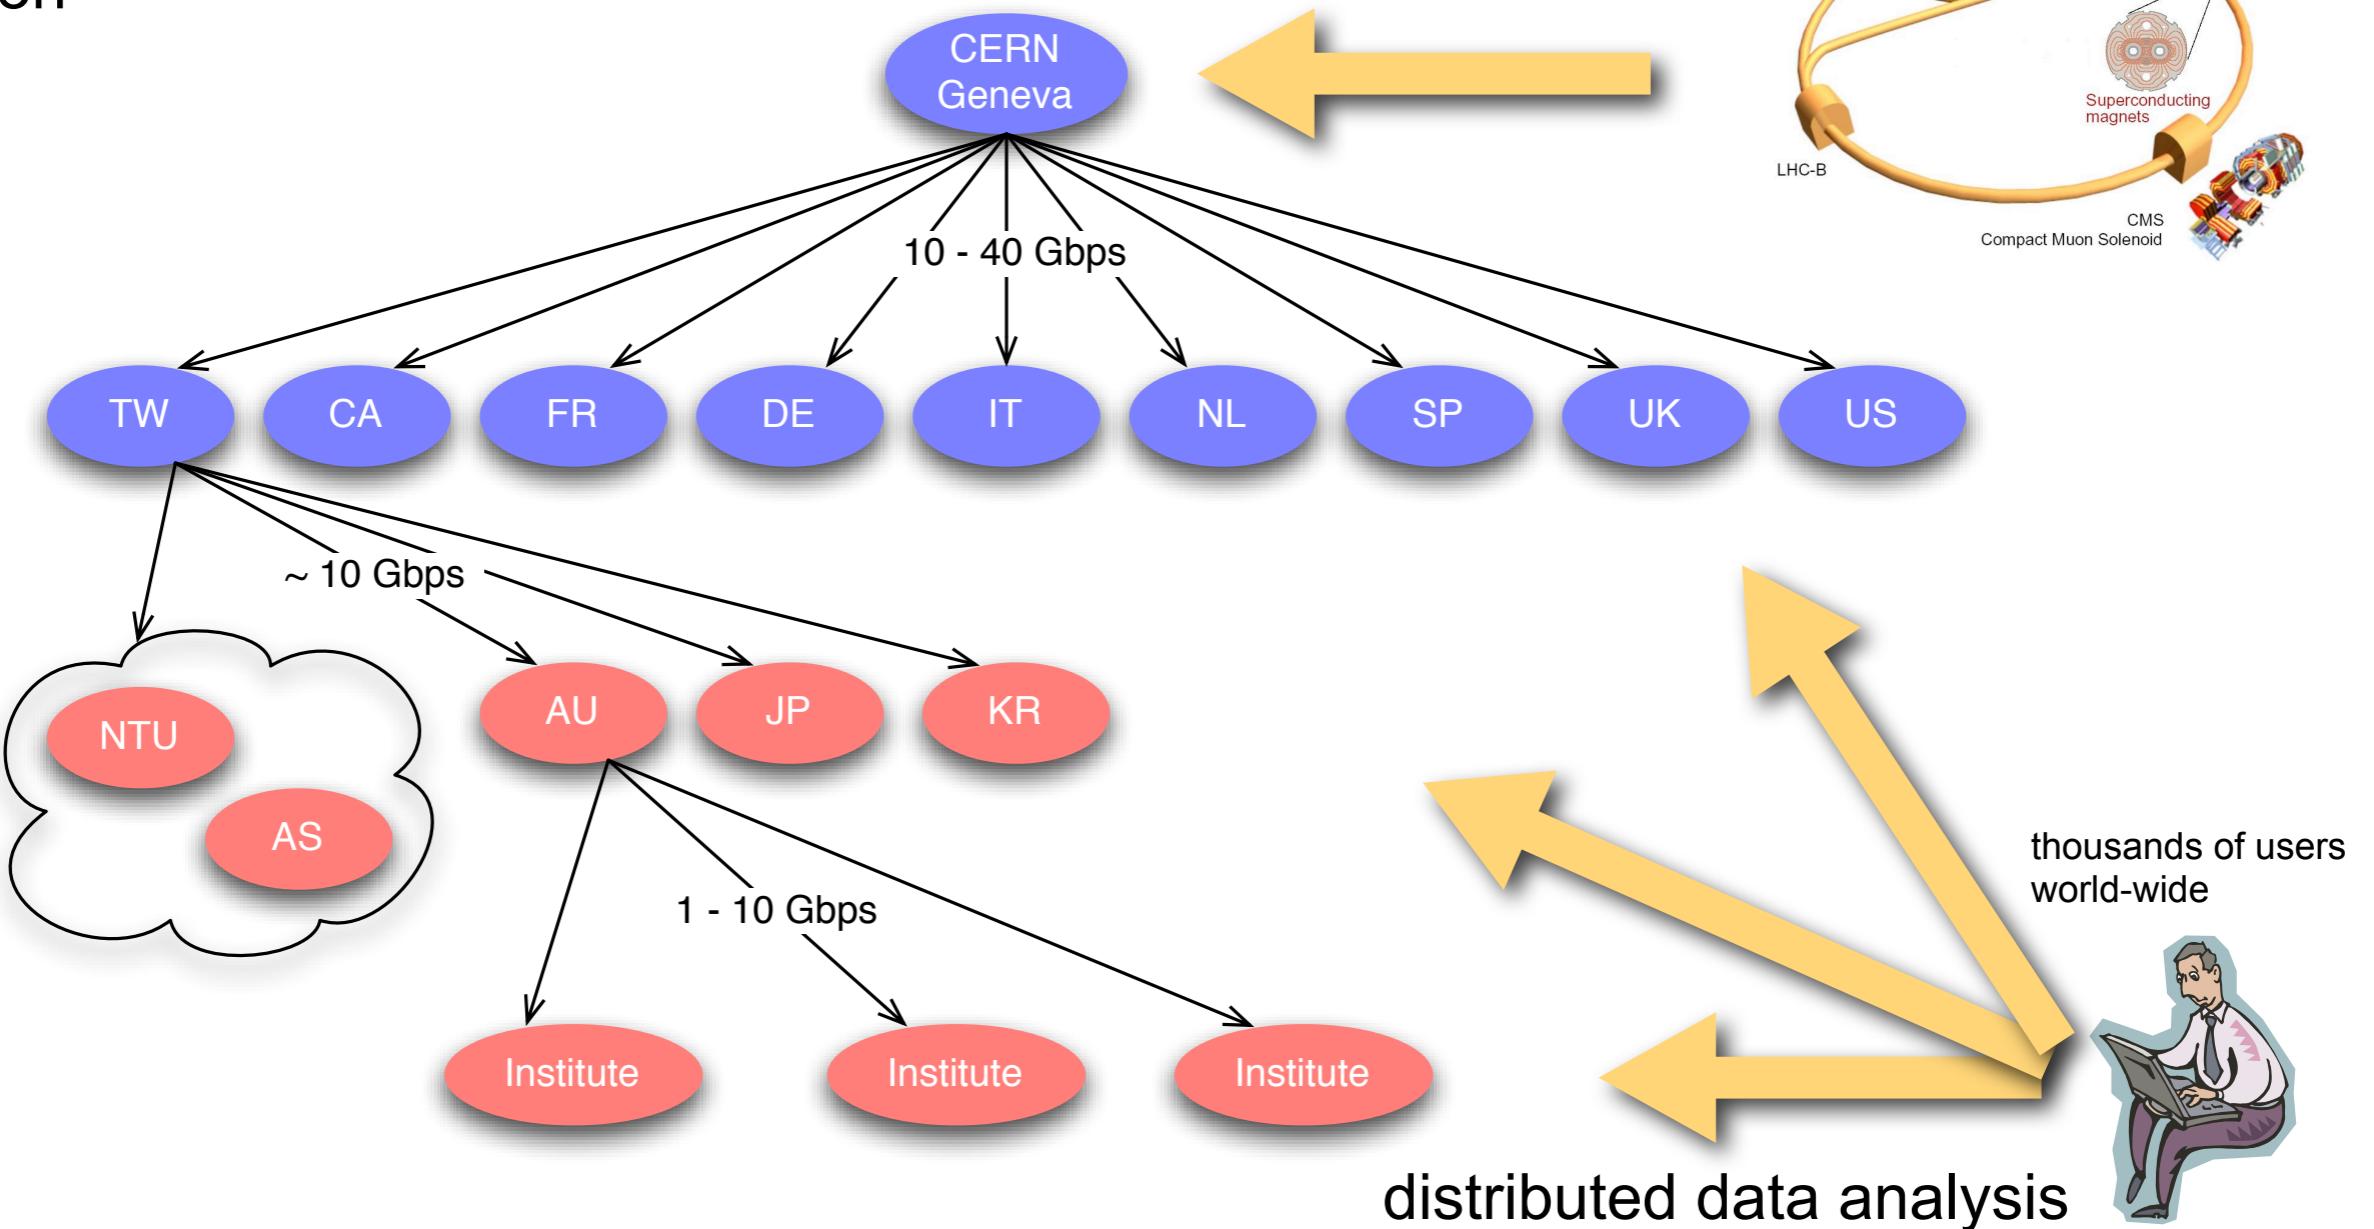
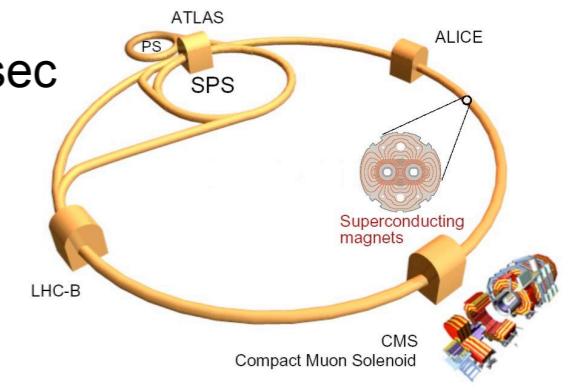
data production rate: $\sim 10 \text{ PByte/year}$

24 x 7 operation

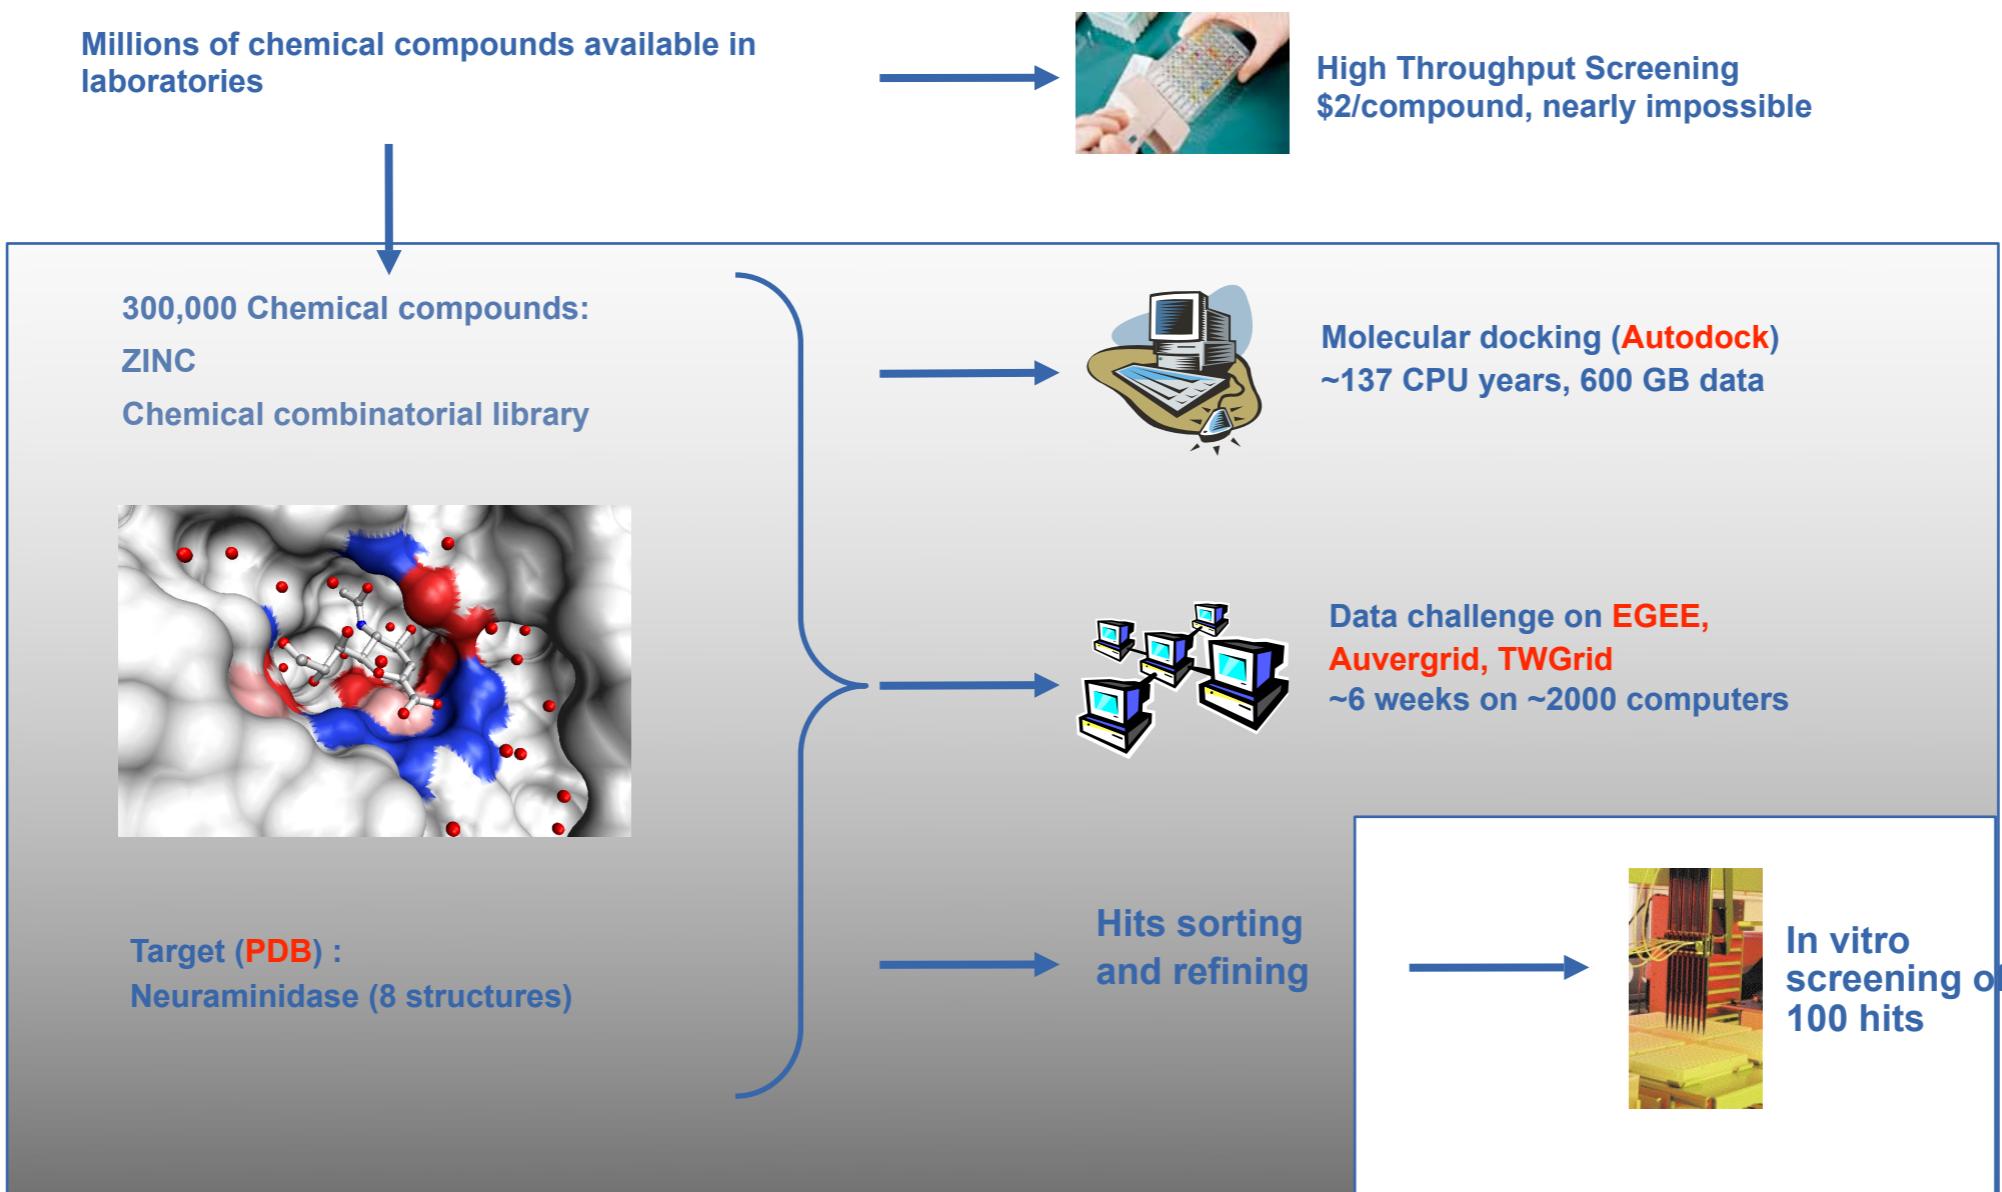
central
production

$\sim 200 - 1500 \text{ MByte/sec}$

$\sim \text{PByte/sec}$

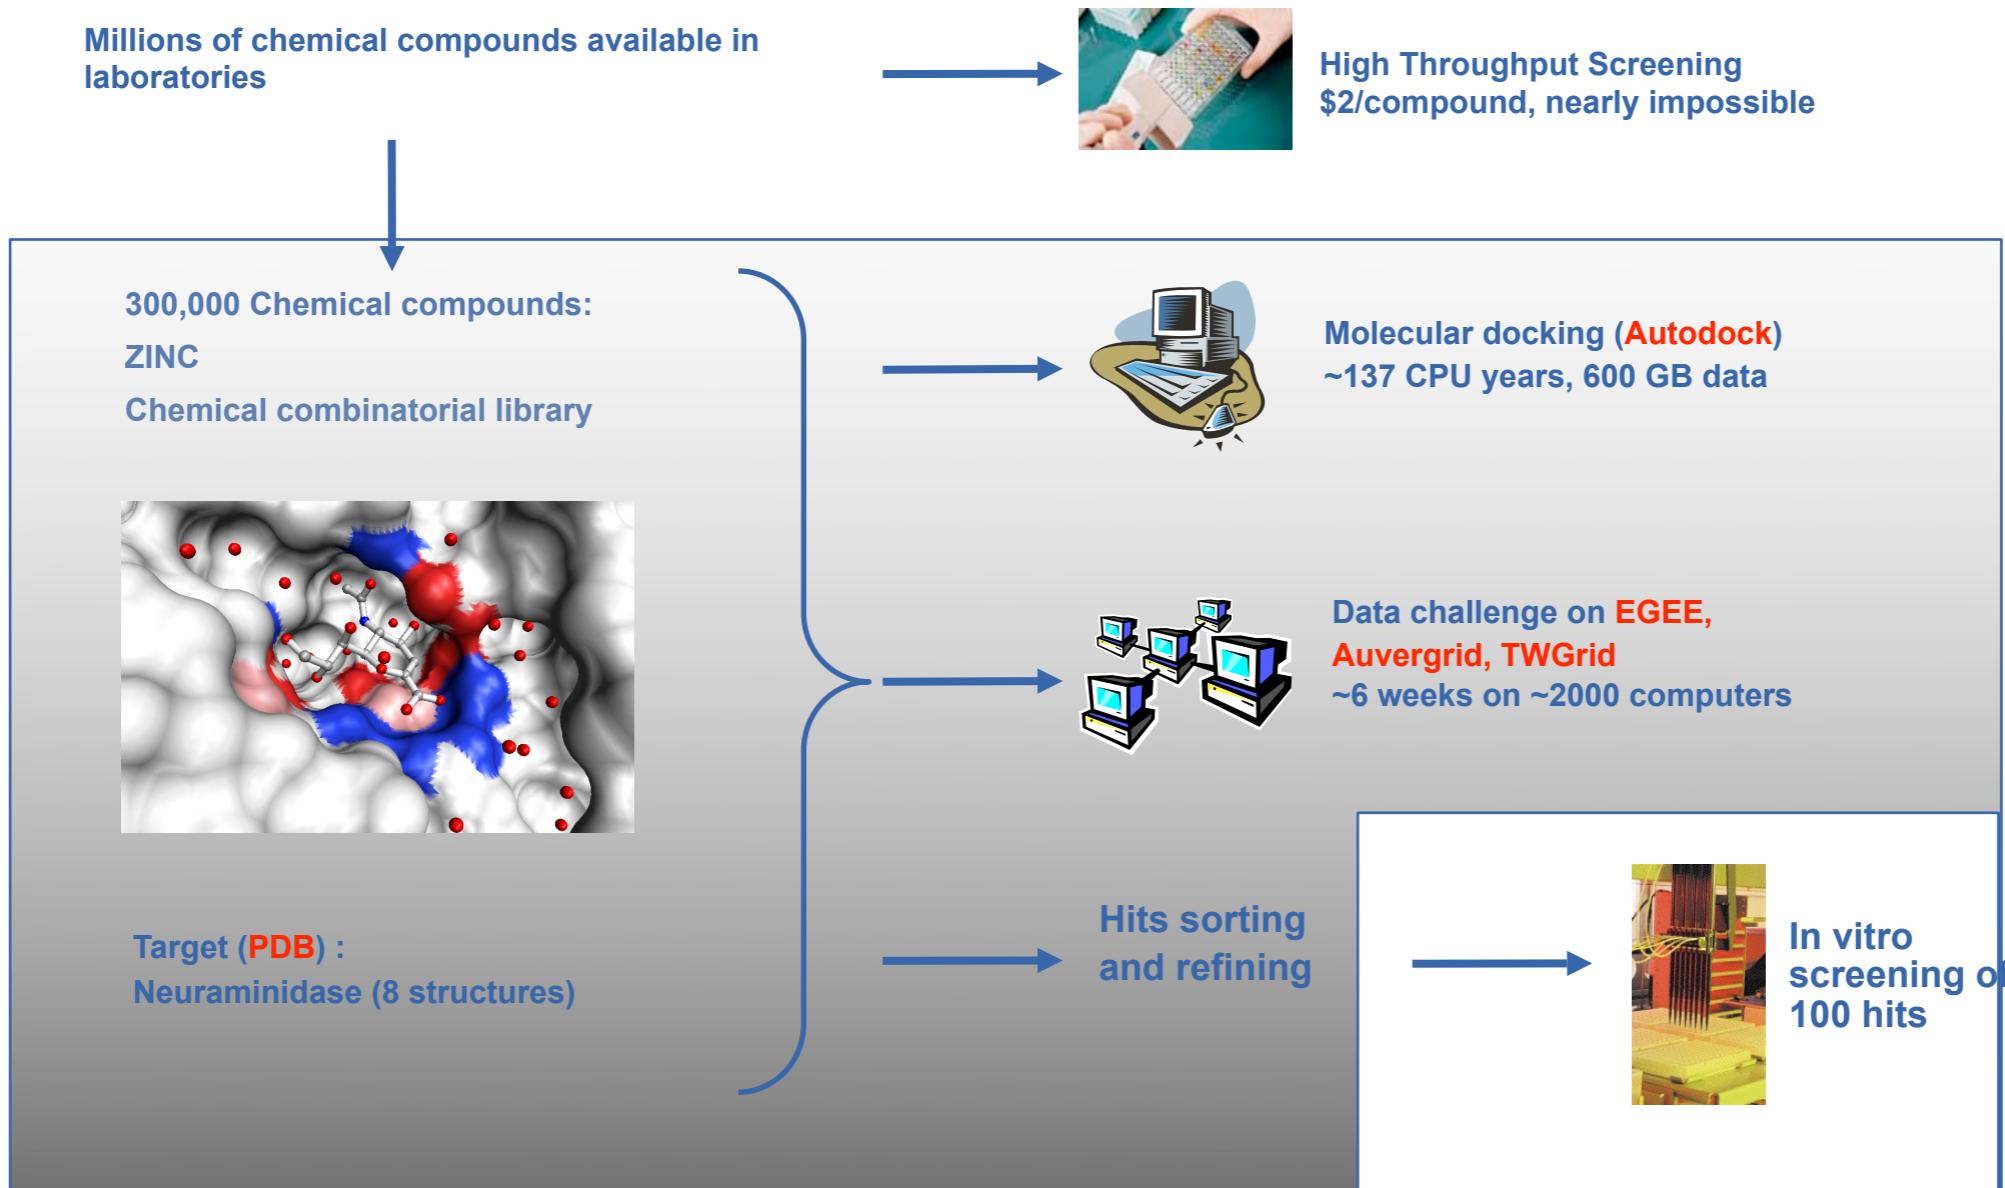


Use case: drug discovery



Use case: drug discovery

on-demand, short deadline, cost effective





The grid software stack

Grid Users

domain experts; non-IT engineers

Grid Applications

domain specific functionalities

Grid Application Utilities

usability, reliability, interoperability

Grid Services (VDT, gLite, ARC, etc.)

monitoring, accounting, brokering

Grid Middleware (Globus)

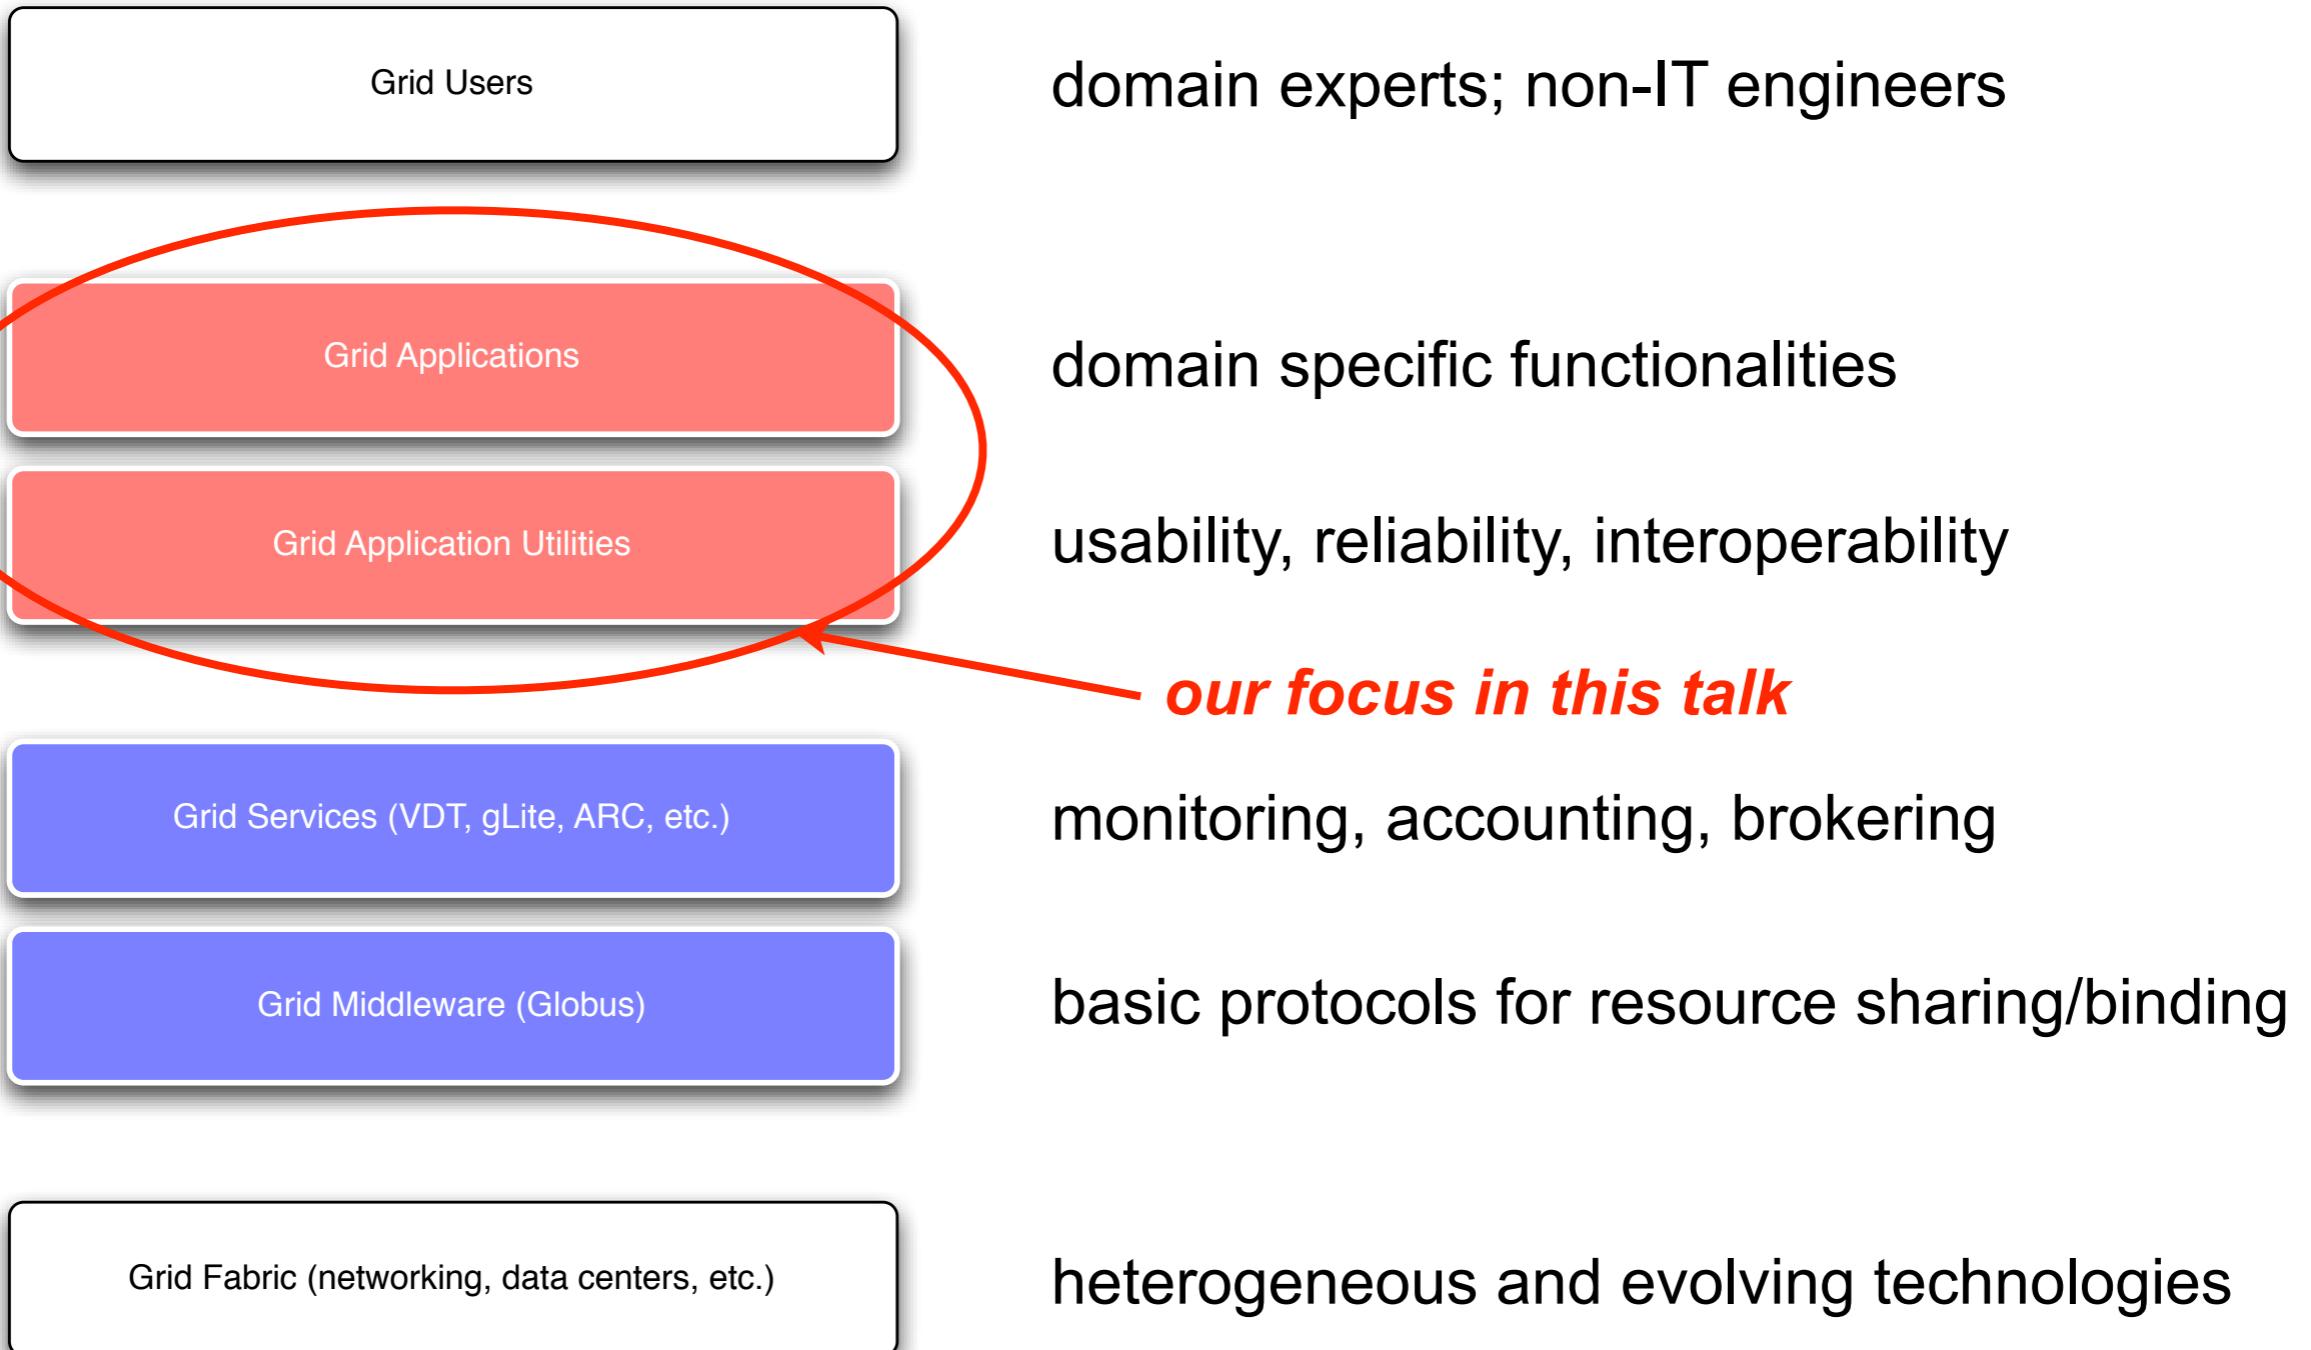
basic protocols for resource sharing/binding

Grid Fabric (networking, data centers, etc.)

heterogeneous and evolving technologies



The grid software stack





Application development issues

- Grid application is a large scale application
 - using thousands of distributed computers to produce millions of files
- Grid is dynamic
 - we have to live with failures
- Grid is generic
 - not well customized for a specific application
- Grid is complex
 - hard to be understood by a normal user



Grid application utilities

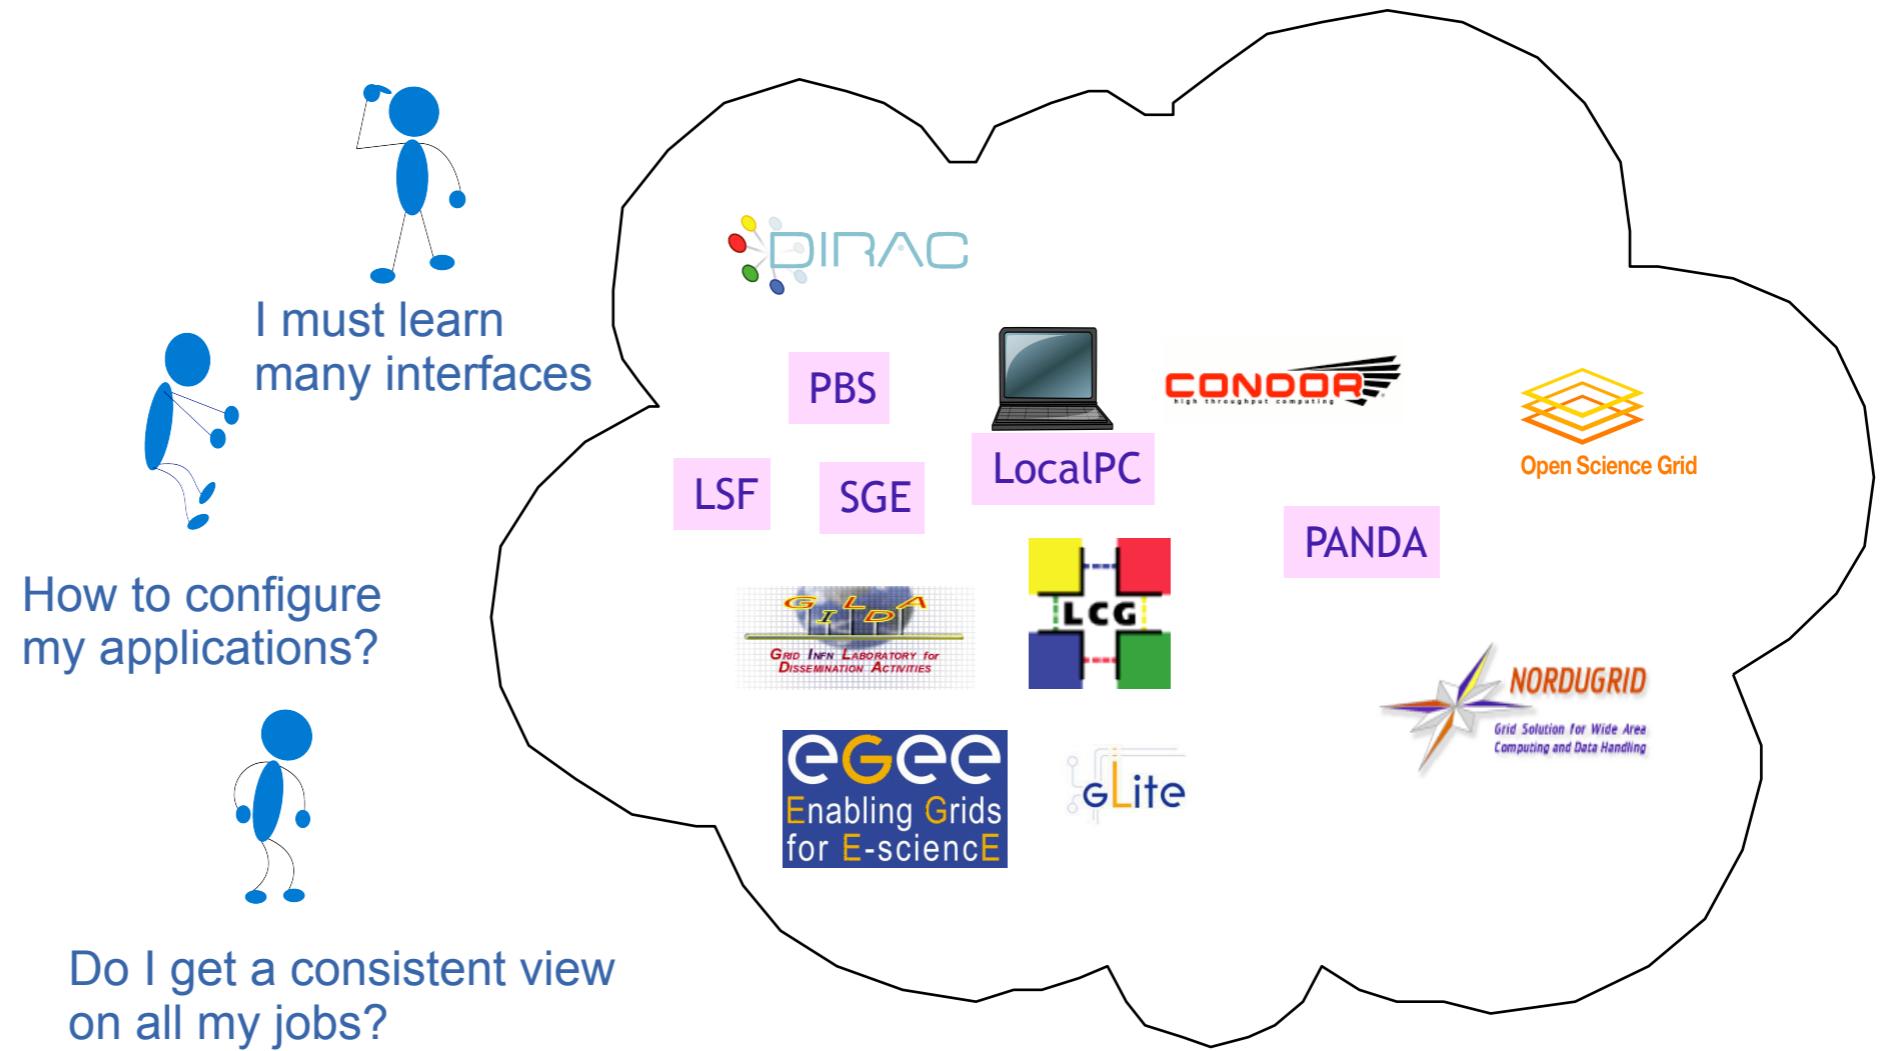
- High-level utilities are needed in building grid applications
 - Also generic; but can be easily customized to meet different application requirements
 - simple, reliable, scalable
 - extensible (i.e. plug-in approach)
- Open source examples
 - GANGA, DIANE, AMGA



GANGA: motivation

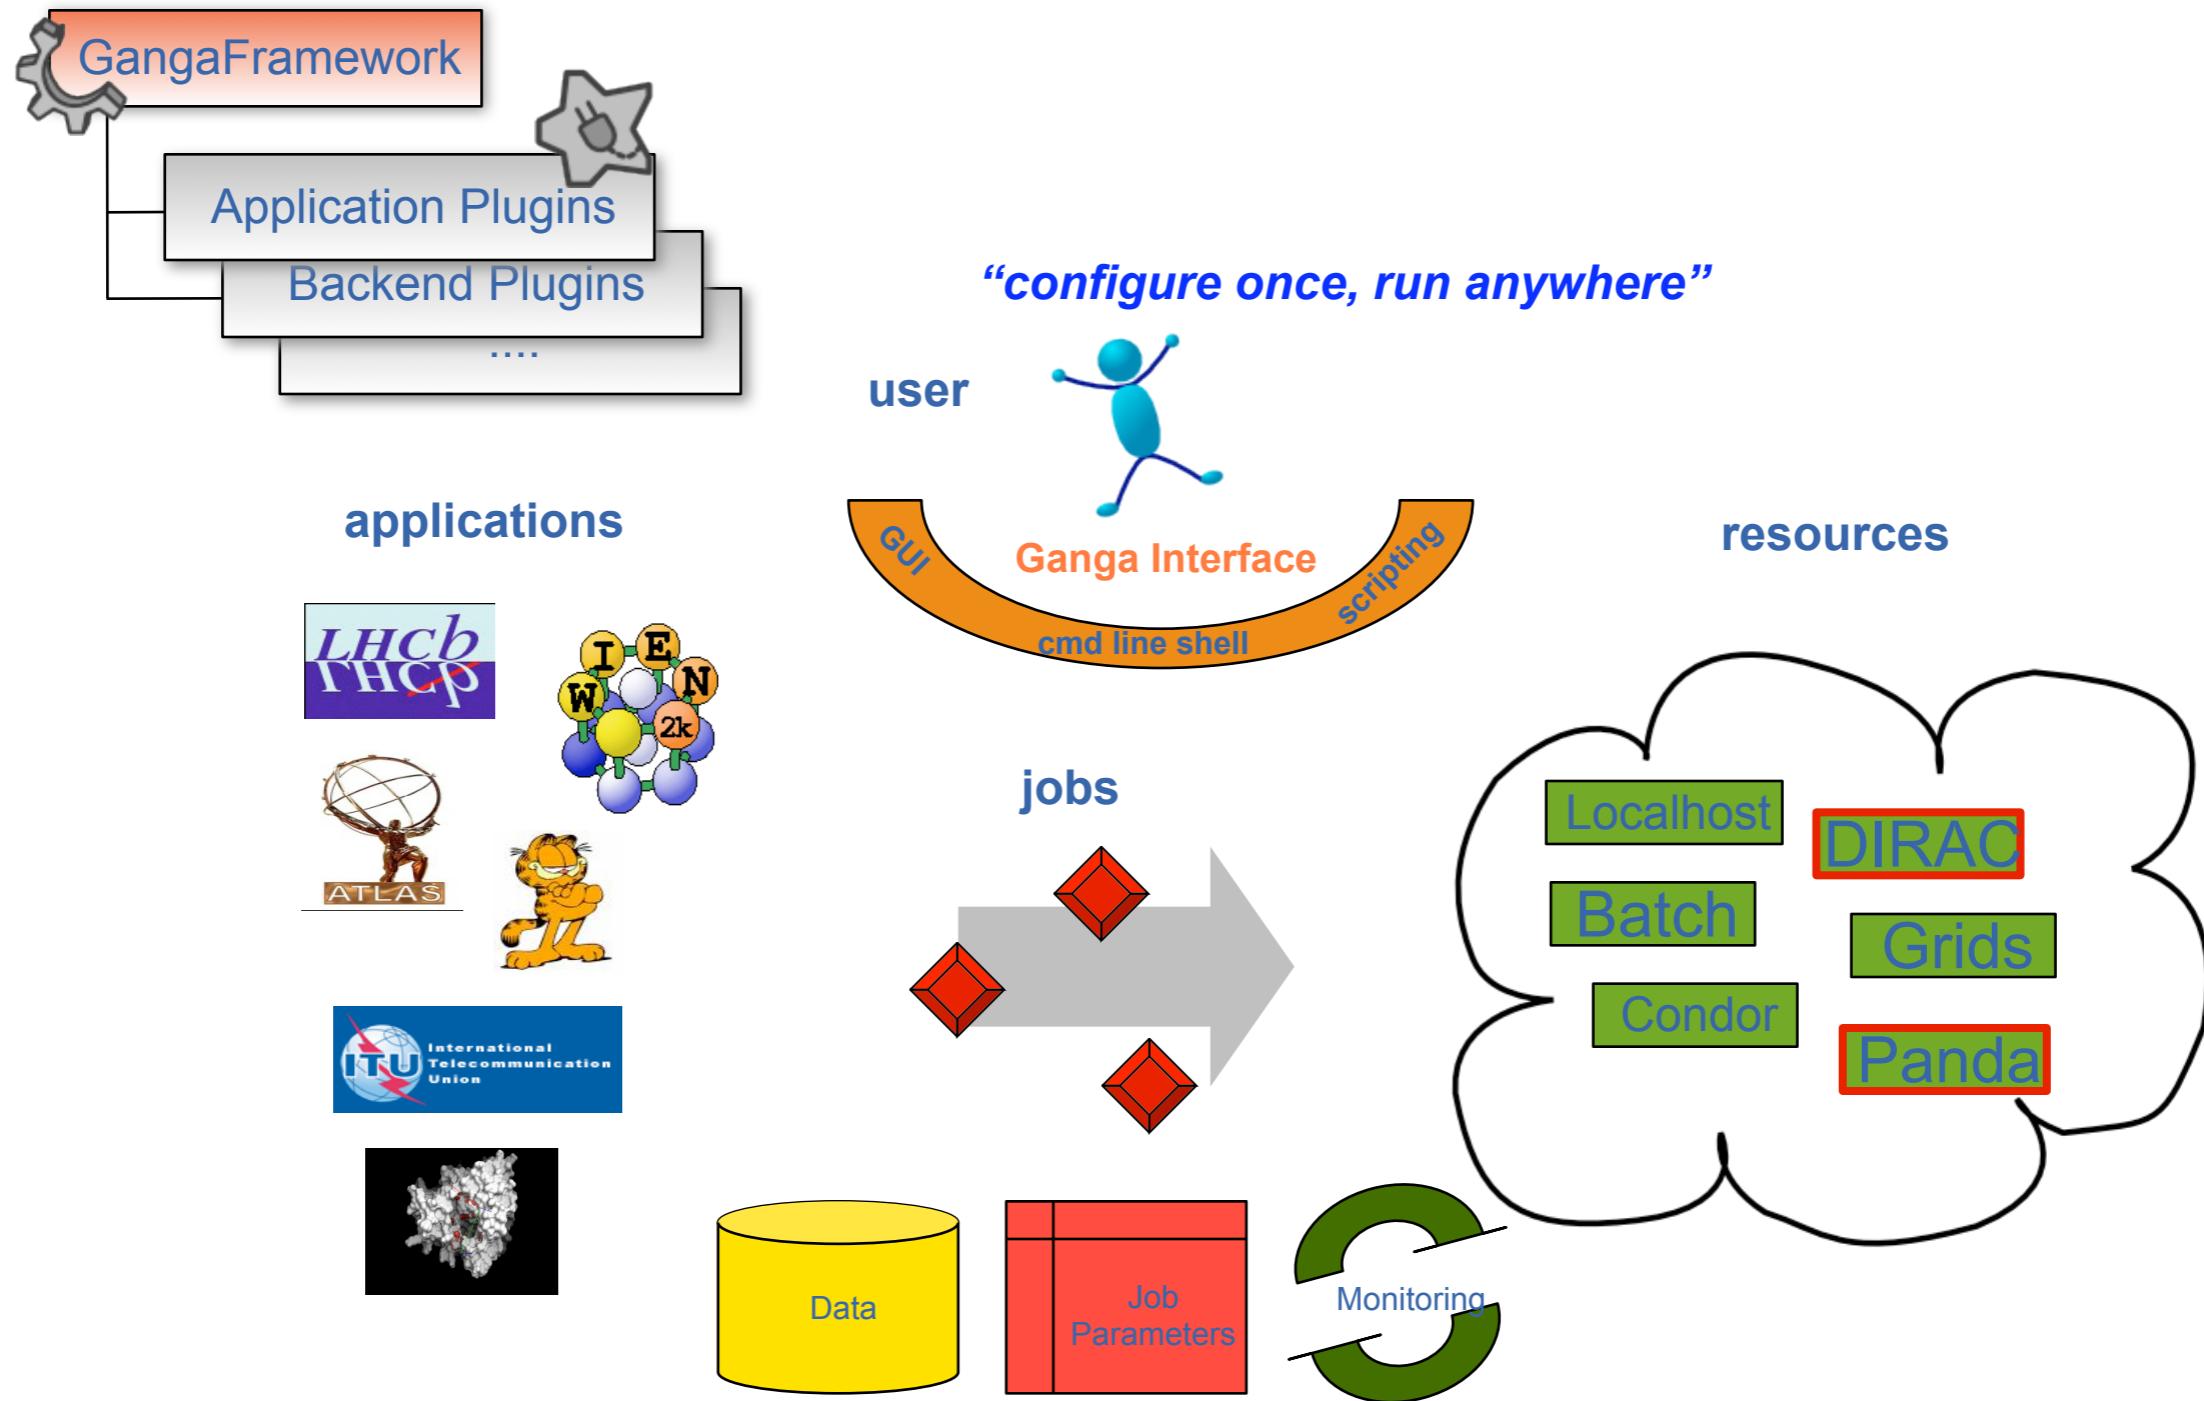


- FAQ: running applications on multiple computing backends
 - grid is a complementary computing environment
 - technology is evolving





GANGA: a high-level job management tool





GANGA: multiple UI flavors

```
*** Welcome to Ganga ***
Version: Ganga-4-2-8
Documentation and support: http://cern.ch/ganga
Type help() or help('index') for online help.

In [1]: jobs
Out[1]: Statistics: 1  jobs
-----
#   id      status      name      subjobs      application
backend.actualCE
#   1    completed

-----
```

CLIP

Executable

The screenshot shows the Ganga GUI window. On the left, there is a table titled "Jobs" listing several job entries. One entry is highlighted in green and labeled "Completed". The table columns include: id, status, name, subjobs, application, exe filename, and backend. The "application" column for the completed job shows "Executable echo". On the right side of the window, there is a panel titled "Job Details" which displays the configuration parameters for the selected job, including its status, name, input and output directories, and application details.

GUI

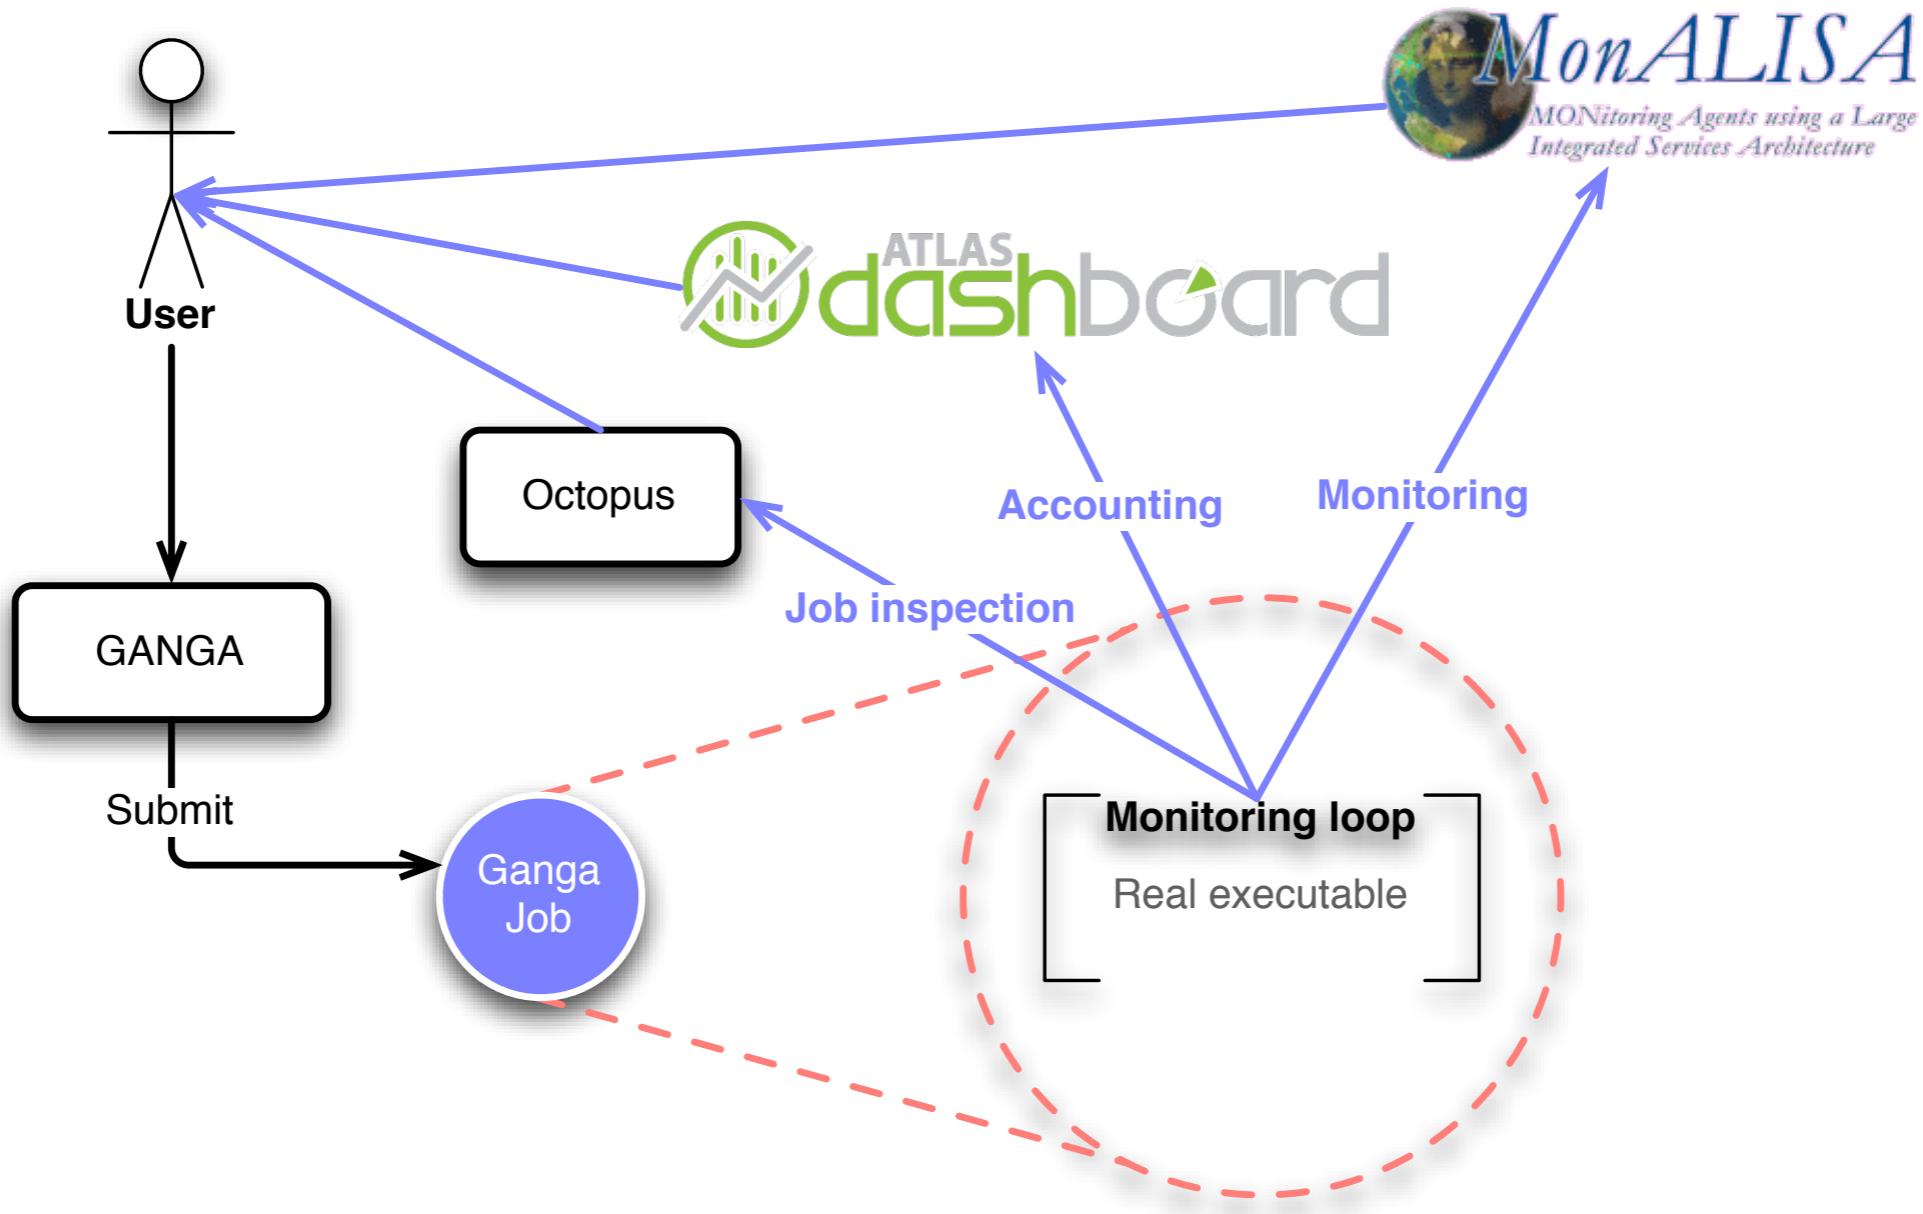
```
#!/usr/bin/env ganga
#-*-python-*-
import time
j = Job()
j.backend = LCG()
j.submit()
while not j.status in ['completed', 'failed']:
    print('job still running')
    time.sleep(30)
```



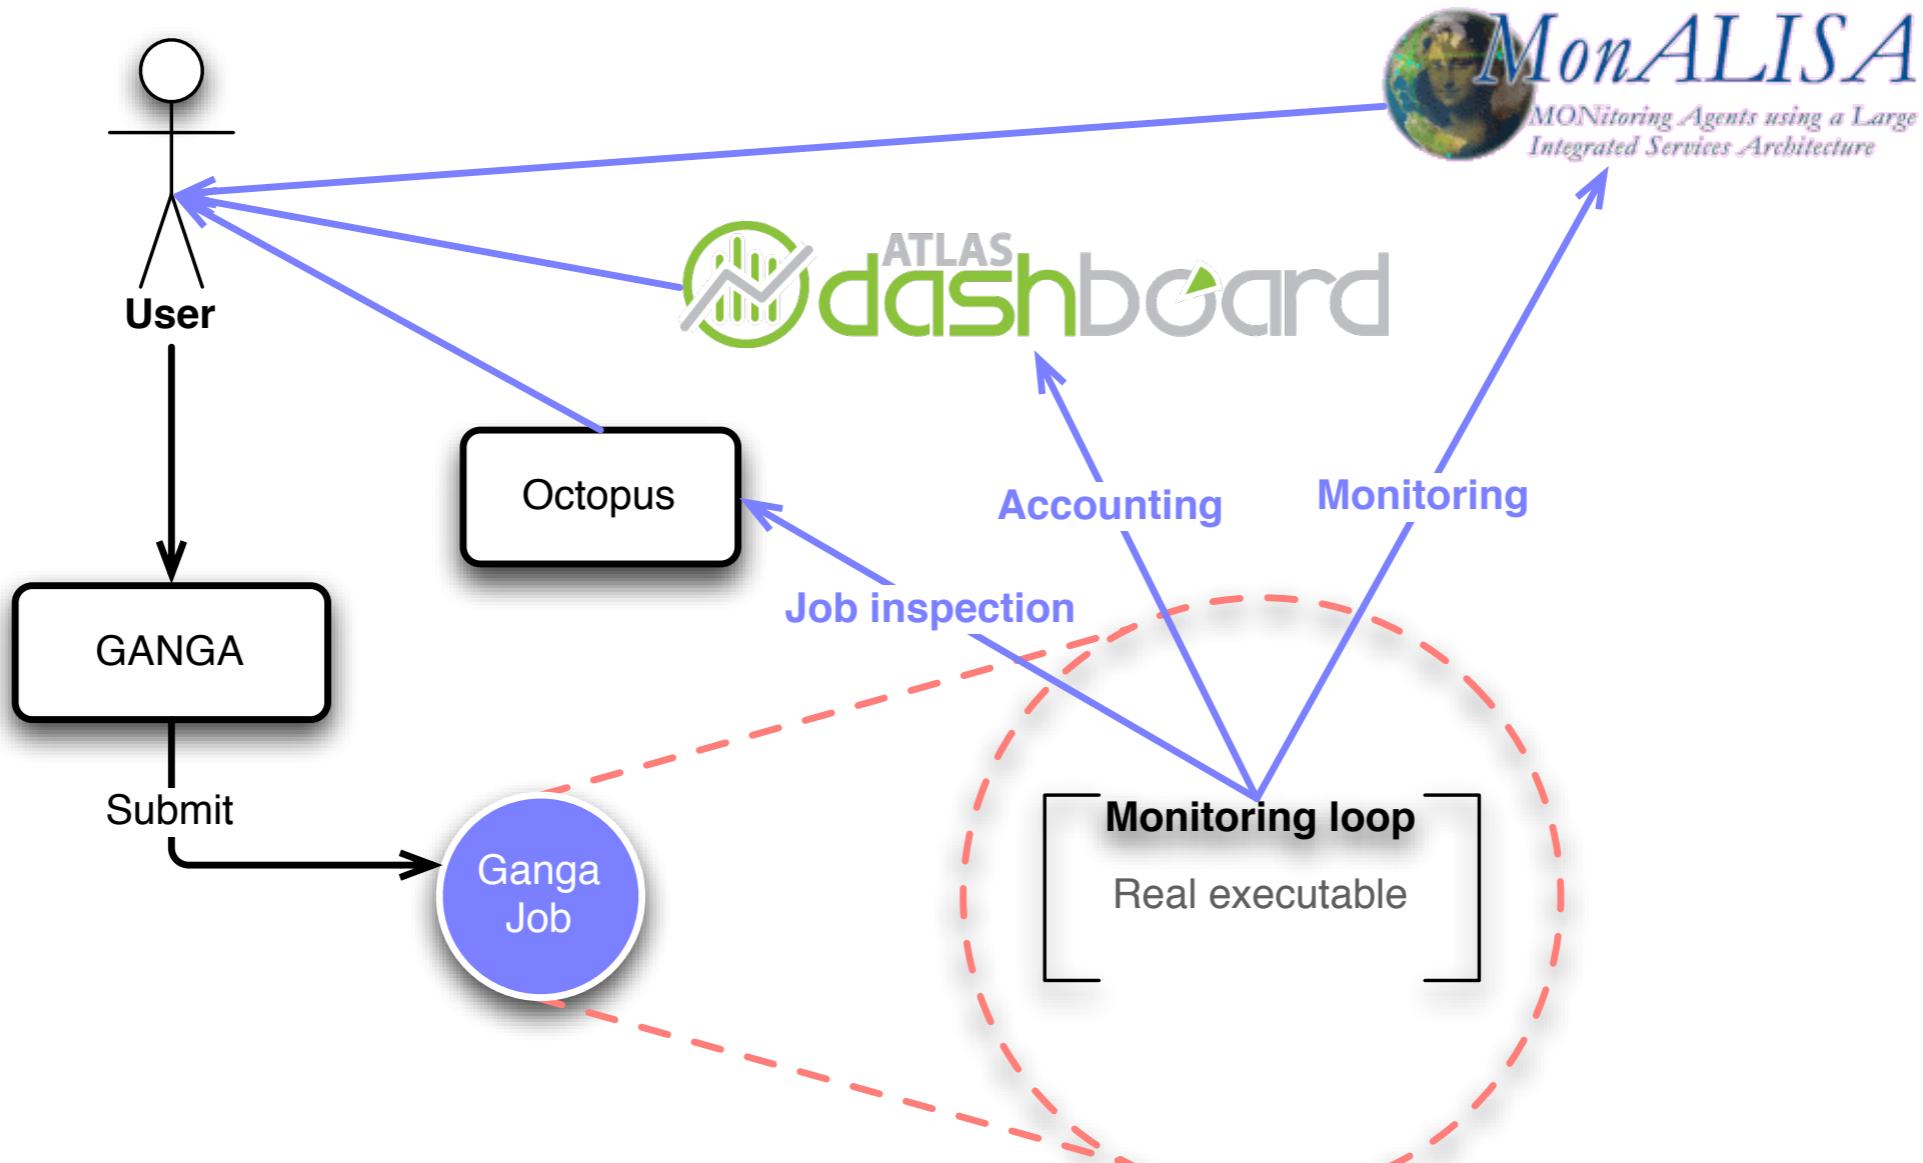
```
./myjob.exec
ganga ./myjob.exec
In [1]: execfile("myjob.exec")
```

GPI & Scripting

GANGA: monitoring, accounting and debugging



GANGA: monitoring, accounting and debugging



modular approach for application specific information

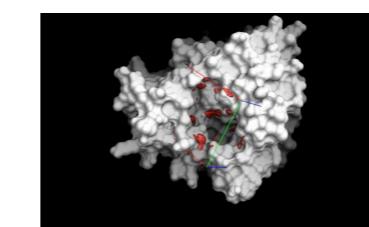


GANGA: users

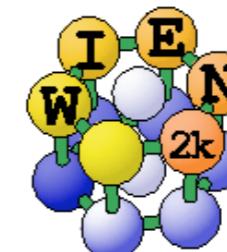


Geant 4

LHCb
~~FHCp~~



Academia Sinica
Genomics Research Center

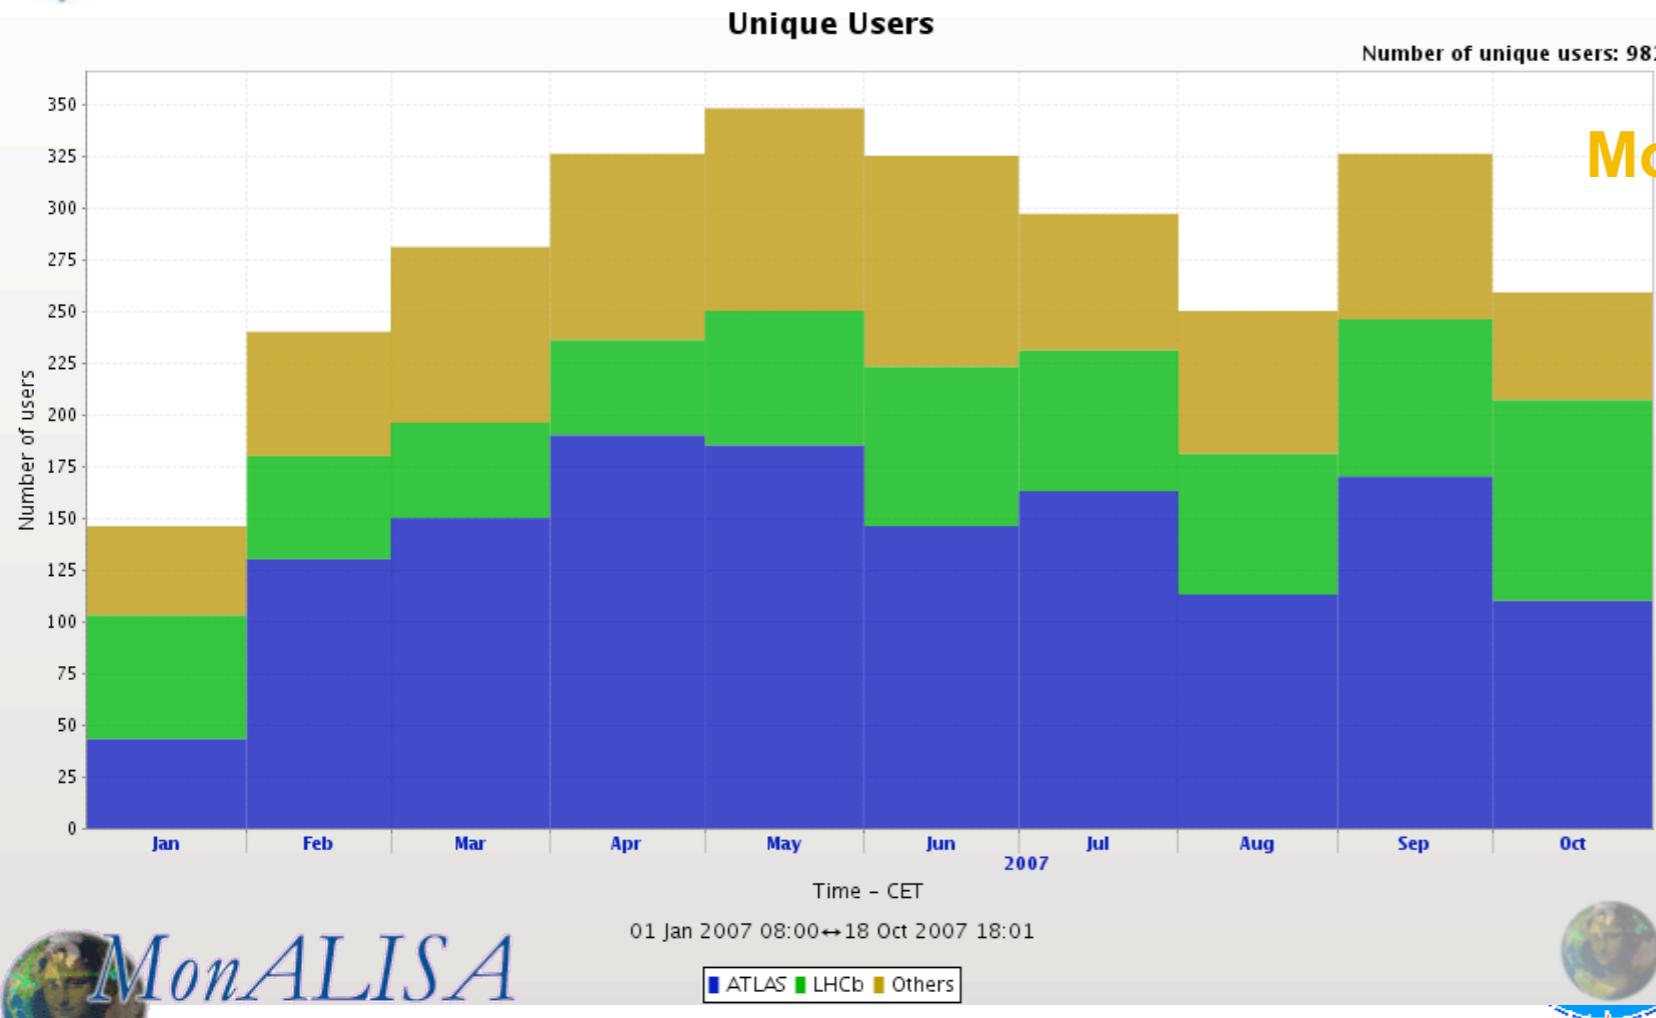


med
austron





GANGA: users



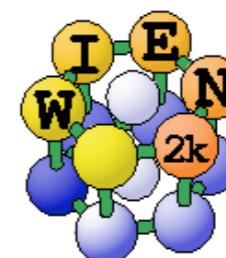
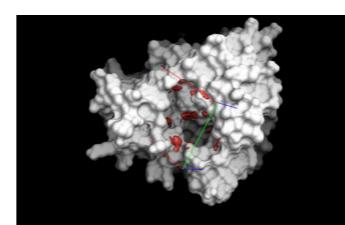
More than 900 users use Ganga



MONitoring Agents using a Large Integrated Services Architecture



Academia Sinica
Genomics Research Center





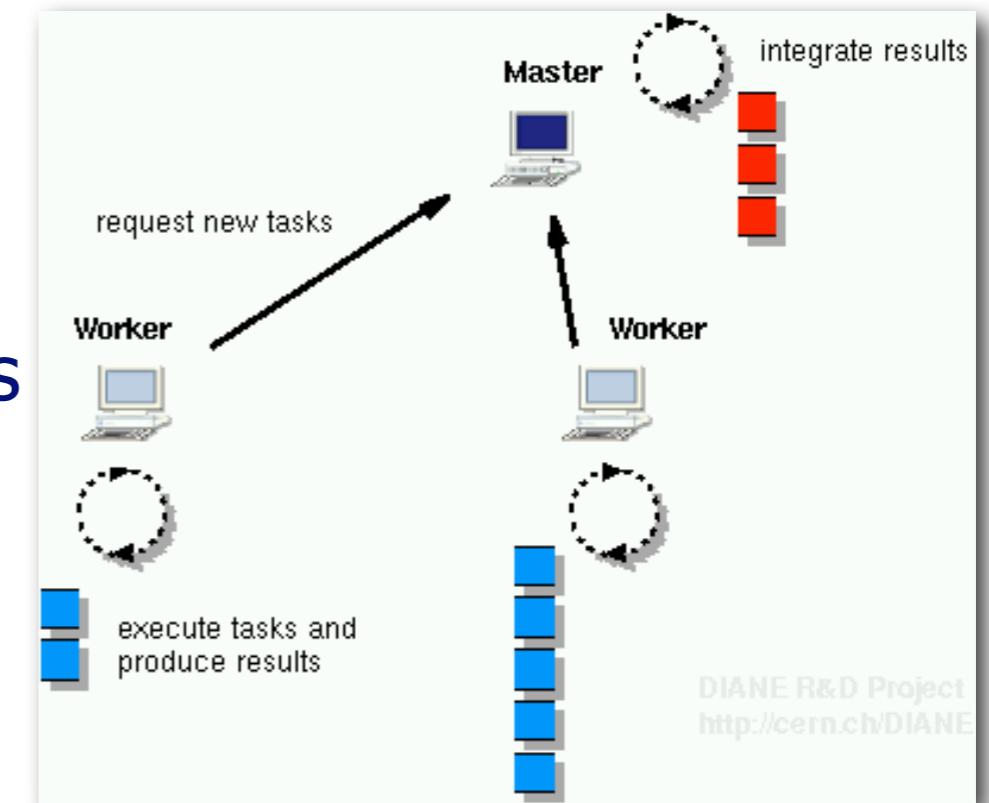
DIANE: motivation



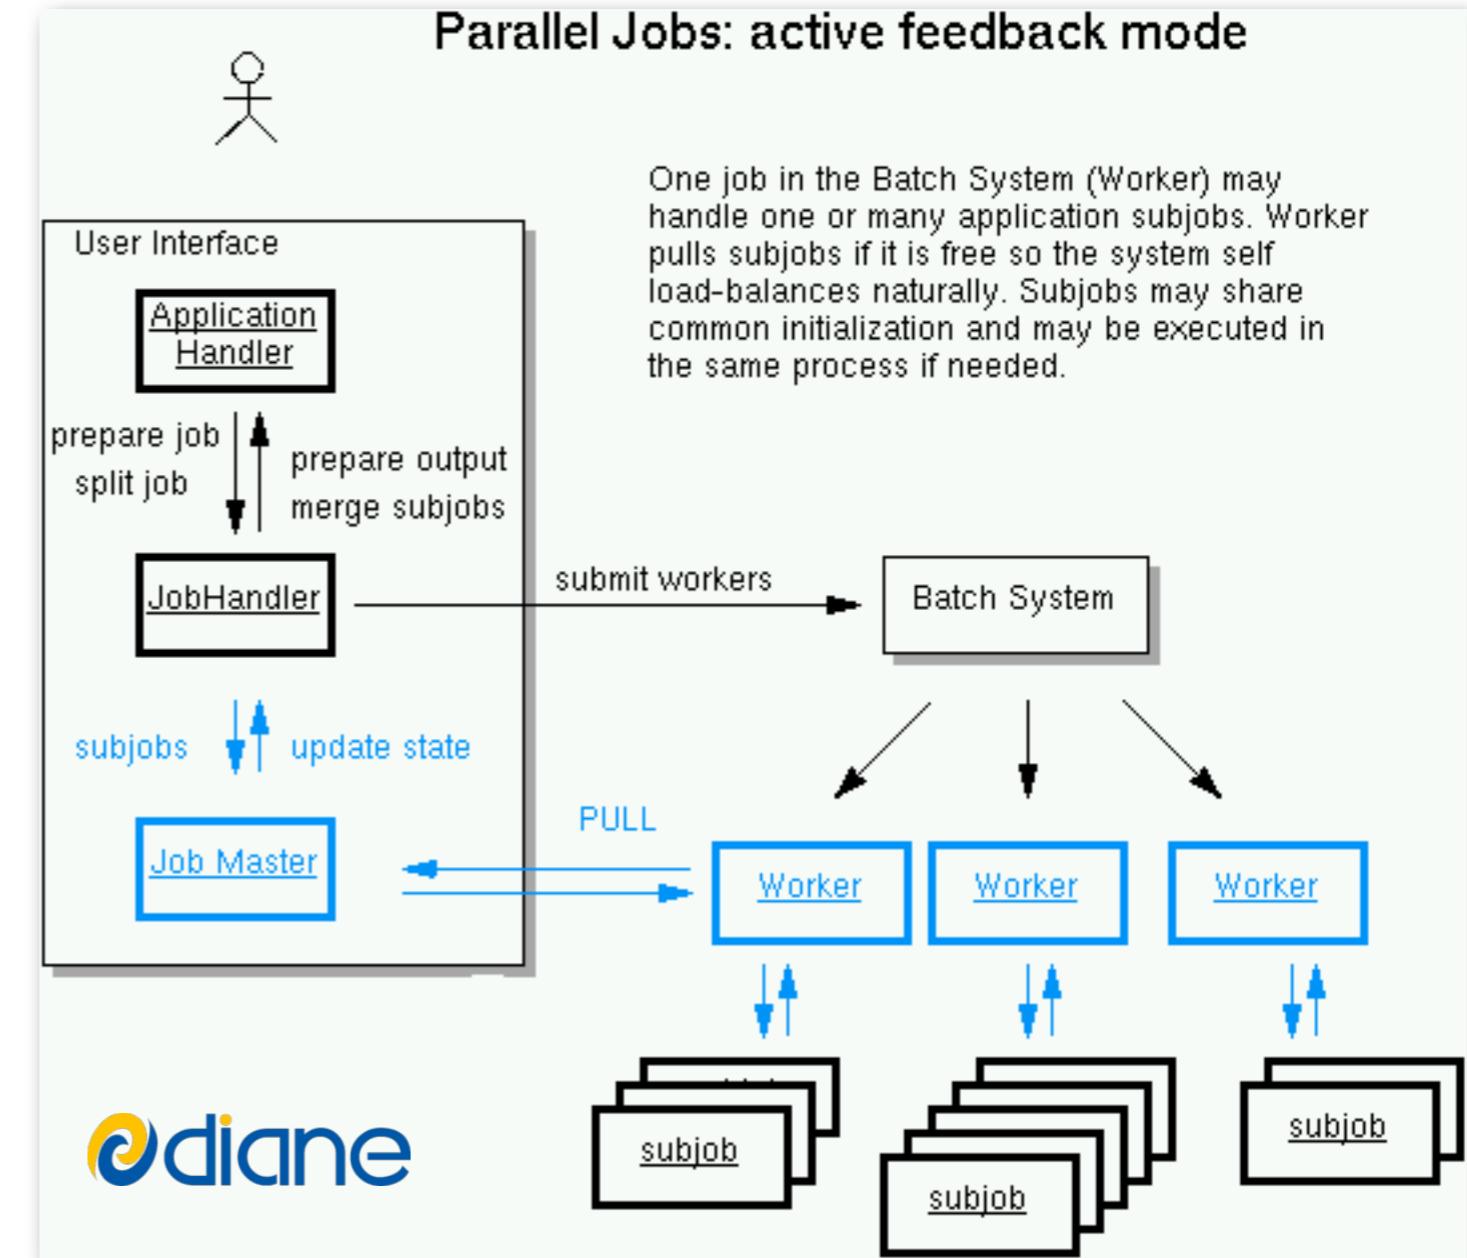
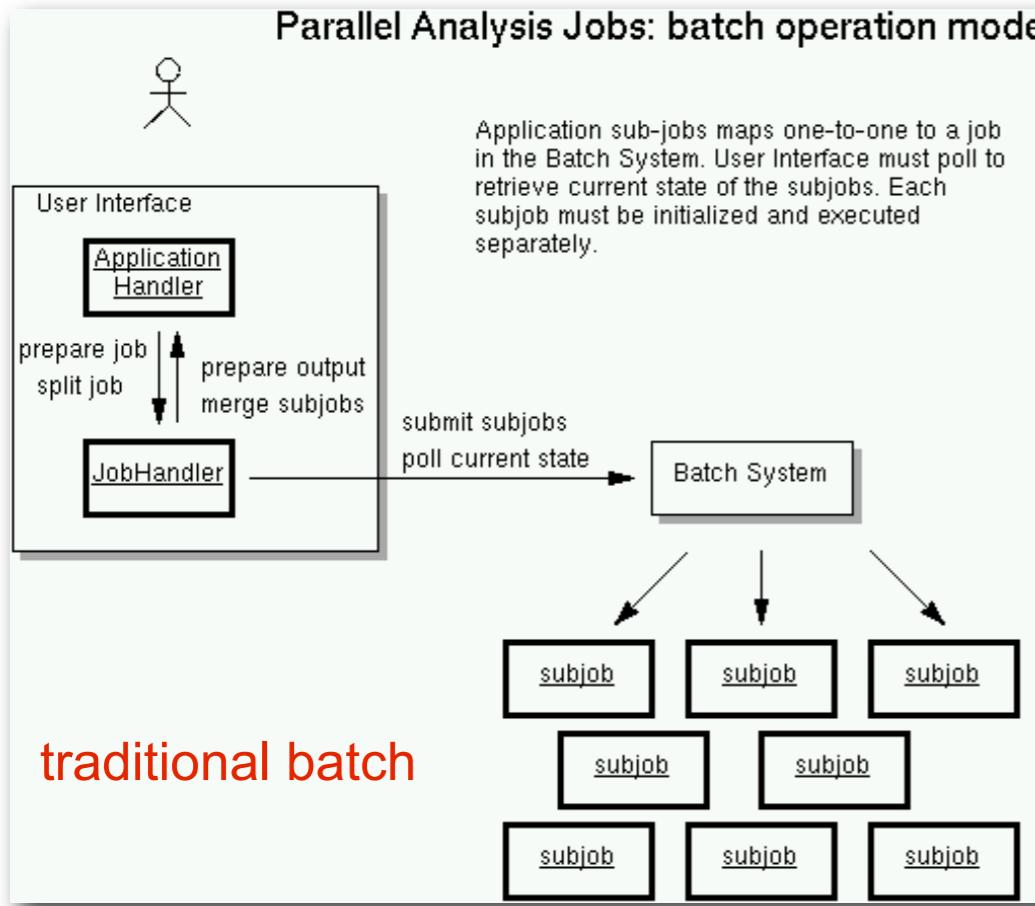
- user(application)-level scheduling
 - advanced resource control by application
 - save the process scheduling overhead
 - efficient computing resource usage
- failure recovery
 - resource error (i.e. network or service unavailable)
 - application error (i.e. library incompatible)
- An use case of GANGA

DIANE: high-level task scheduler

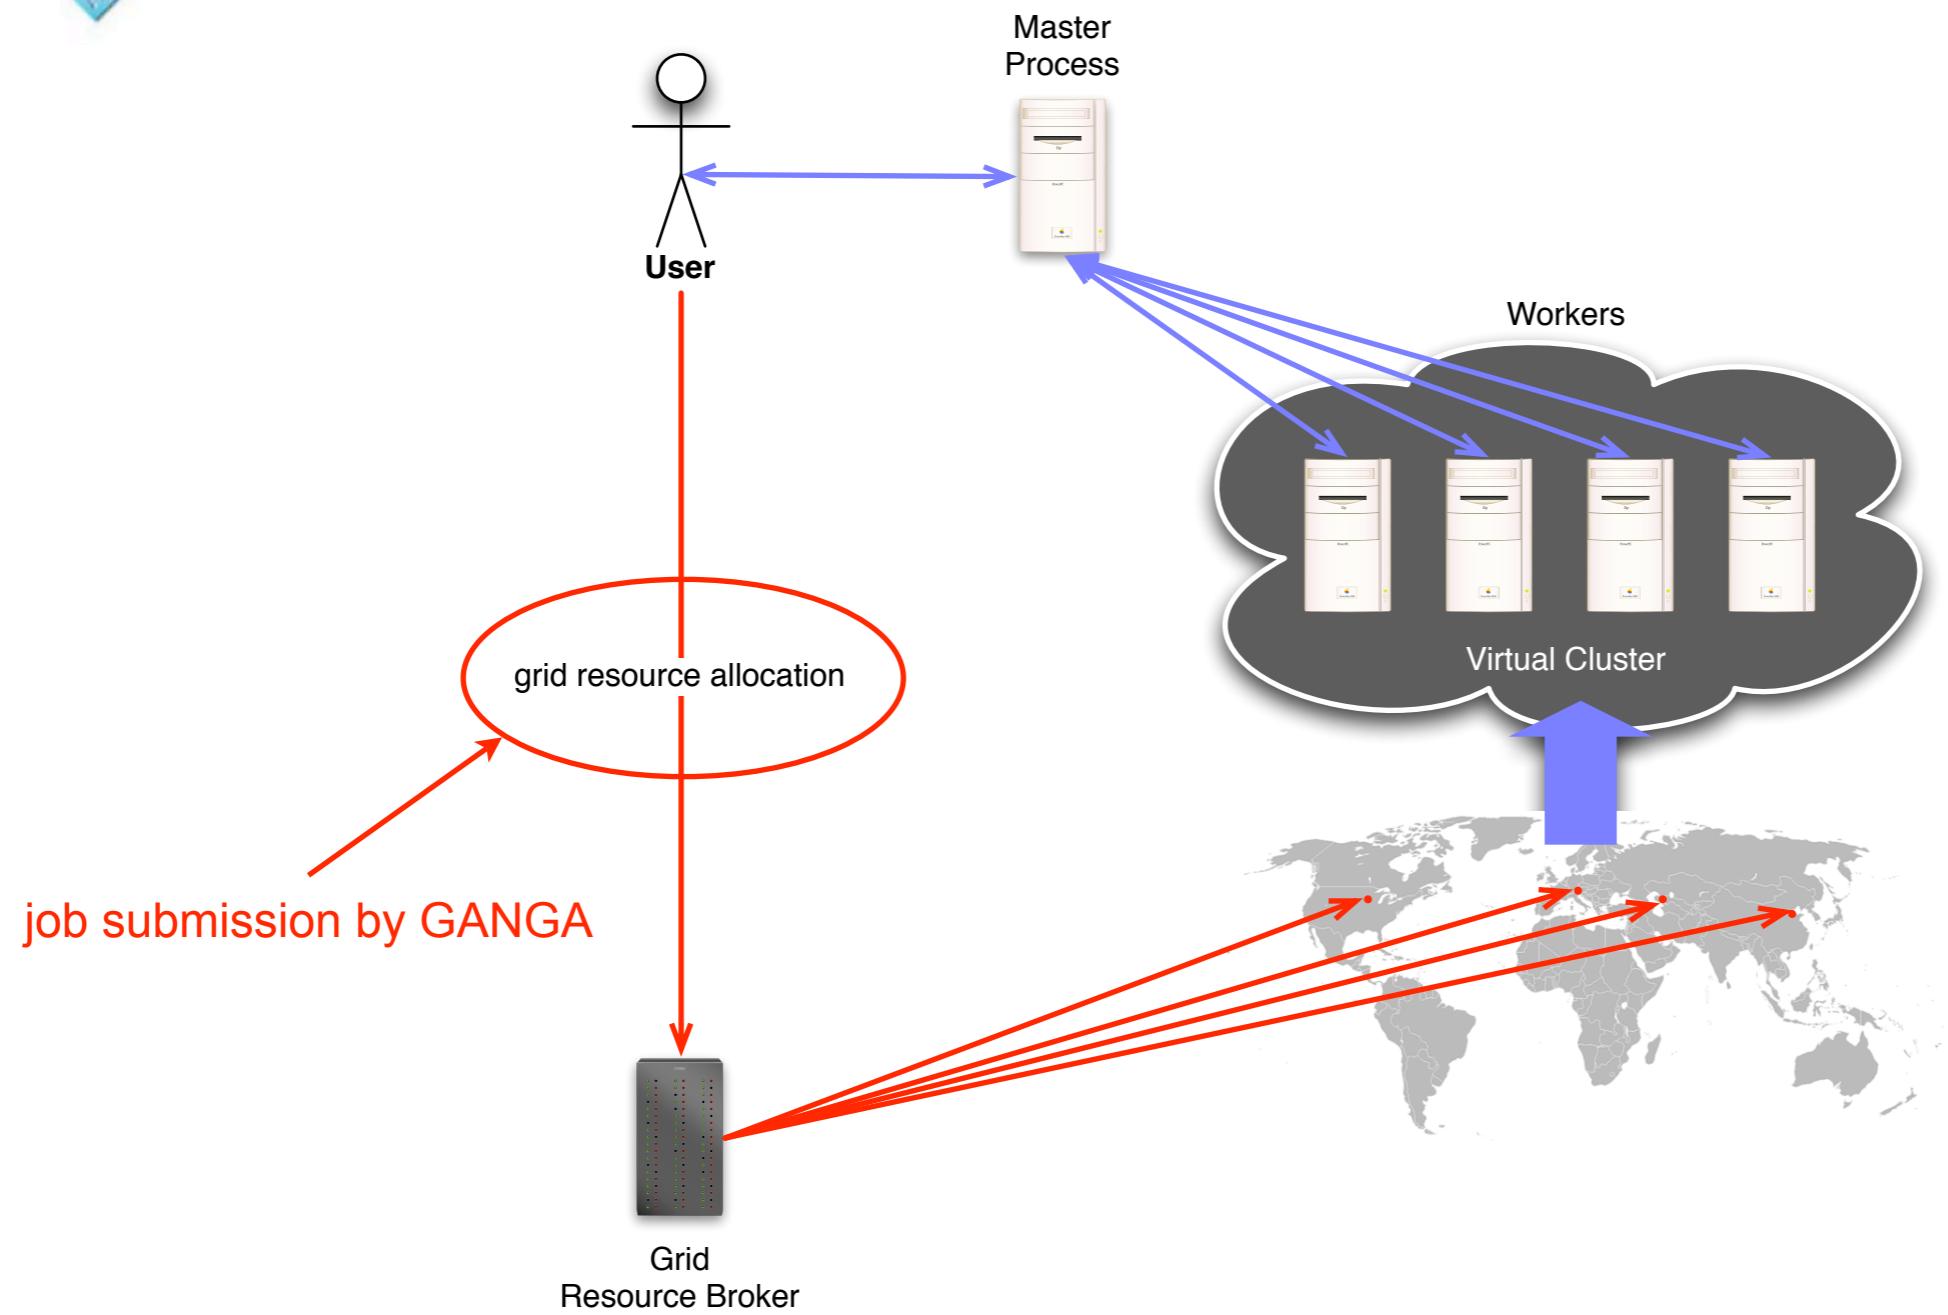
- A lightweight framework for parallel scientific applications in master-worker model
 - assuming that job can be split into a number of independent tasks
- It takes care of all synchronization, communication and workflow management details on behalf of application
 - it's a framework
- Heartbeat worker health check
- Plug-ins for specific applications
- Application defined fail-over mechanisms



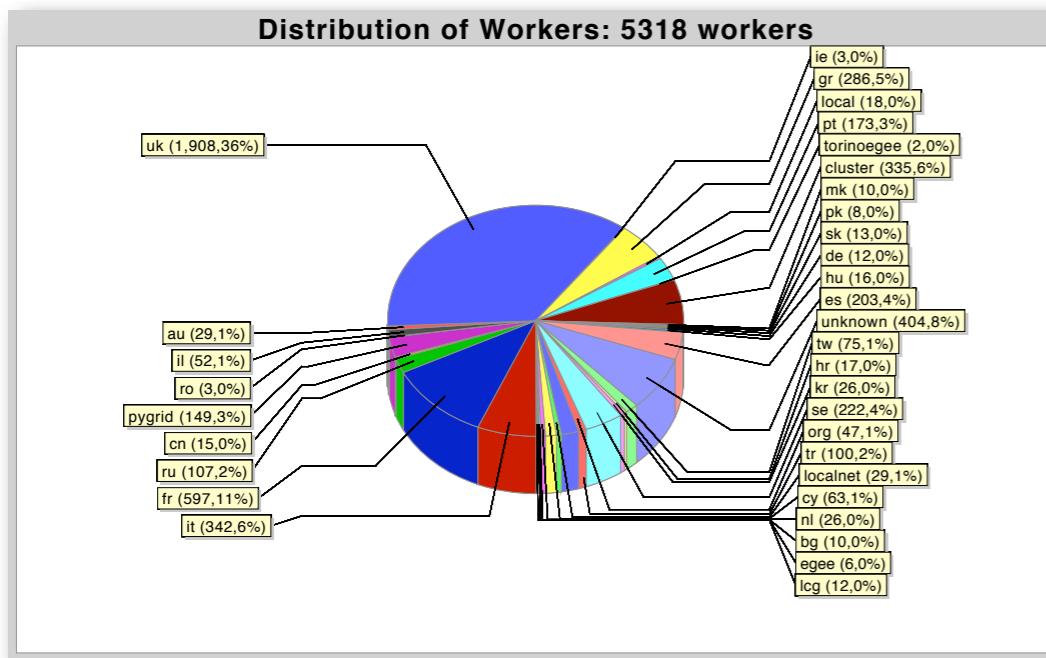
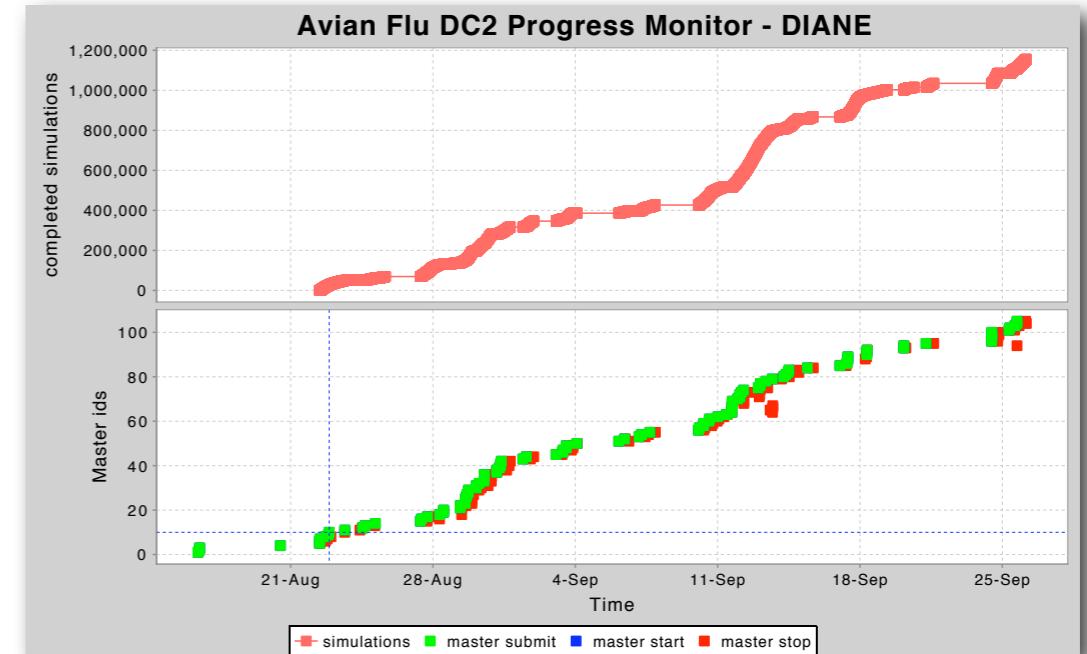
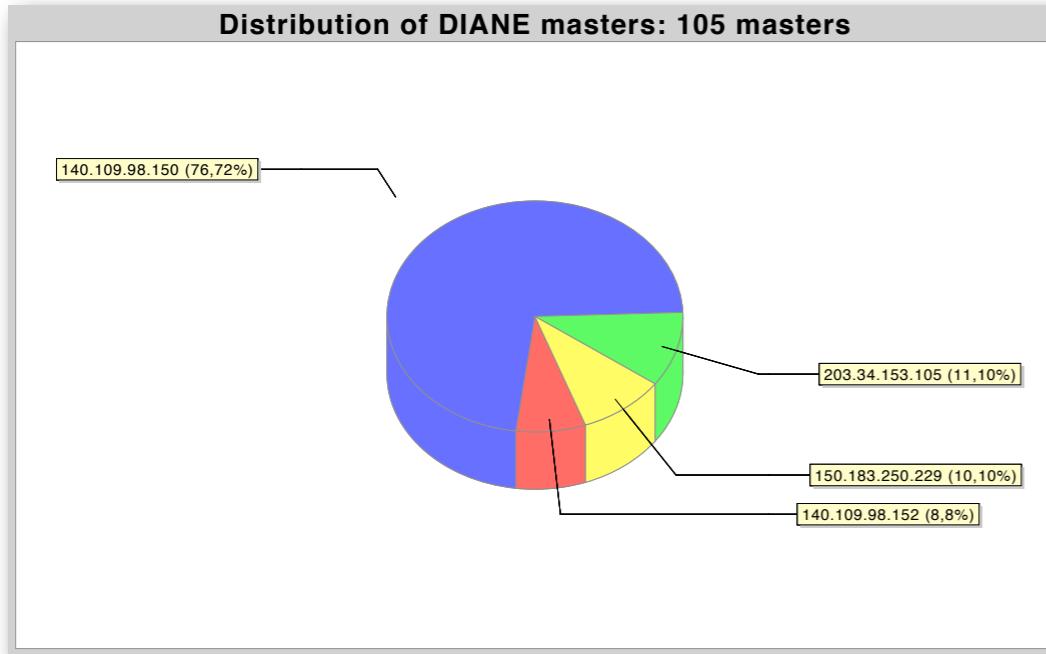
DIANE: job pull model



DIANE: a grid overlay

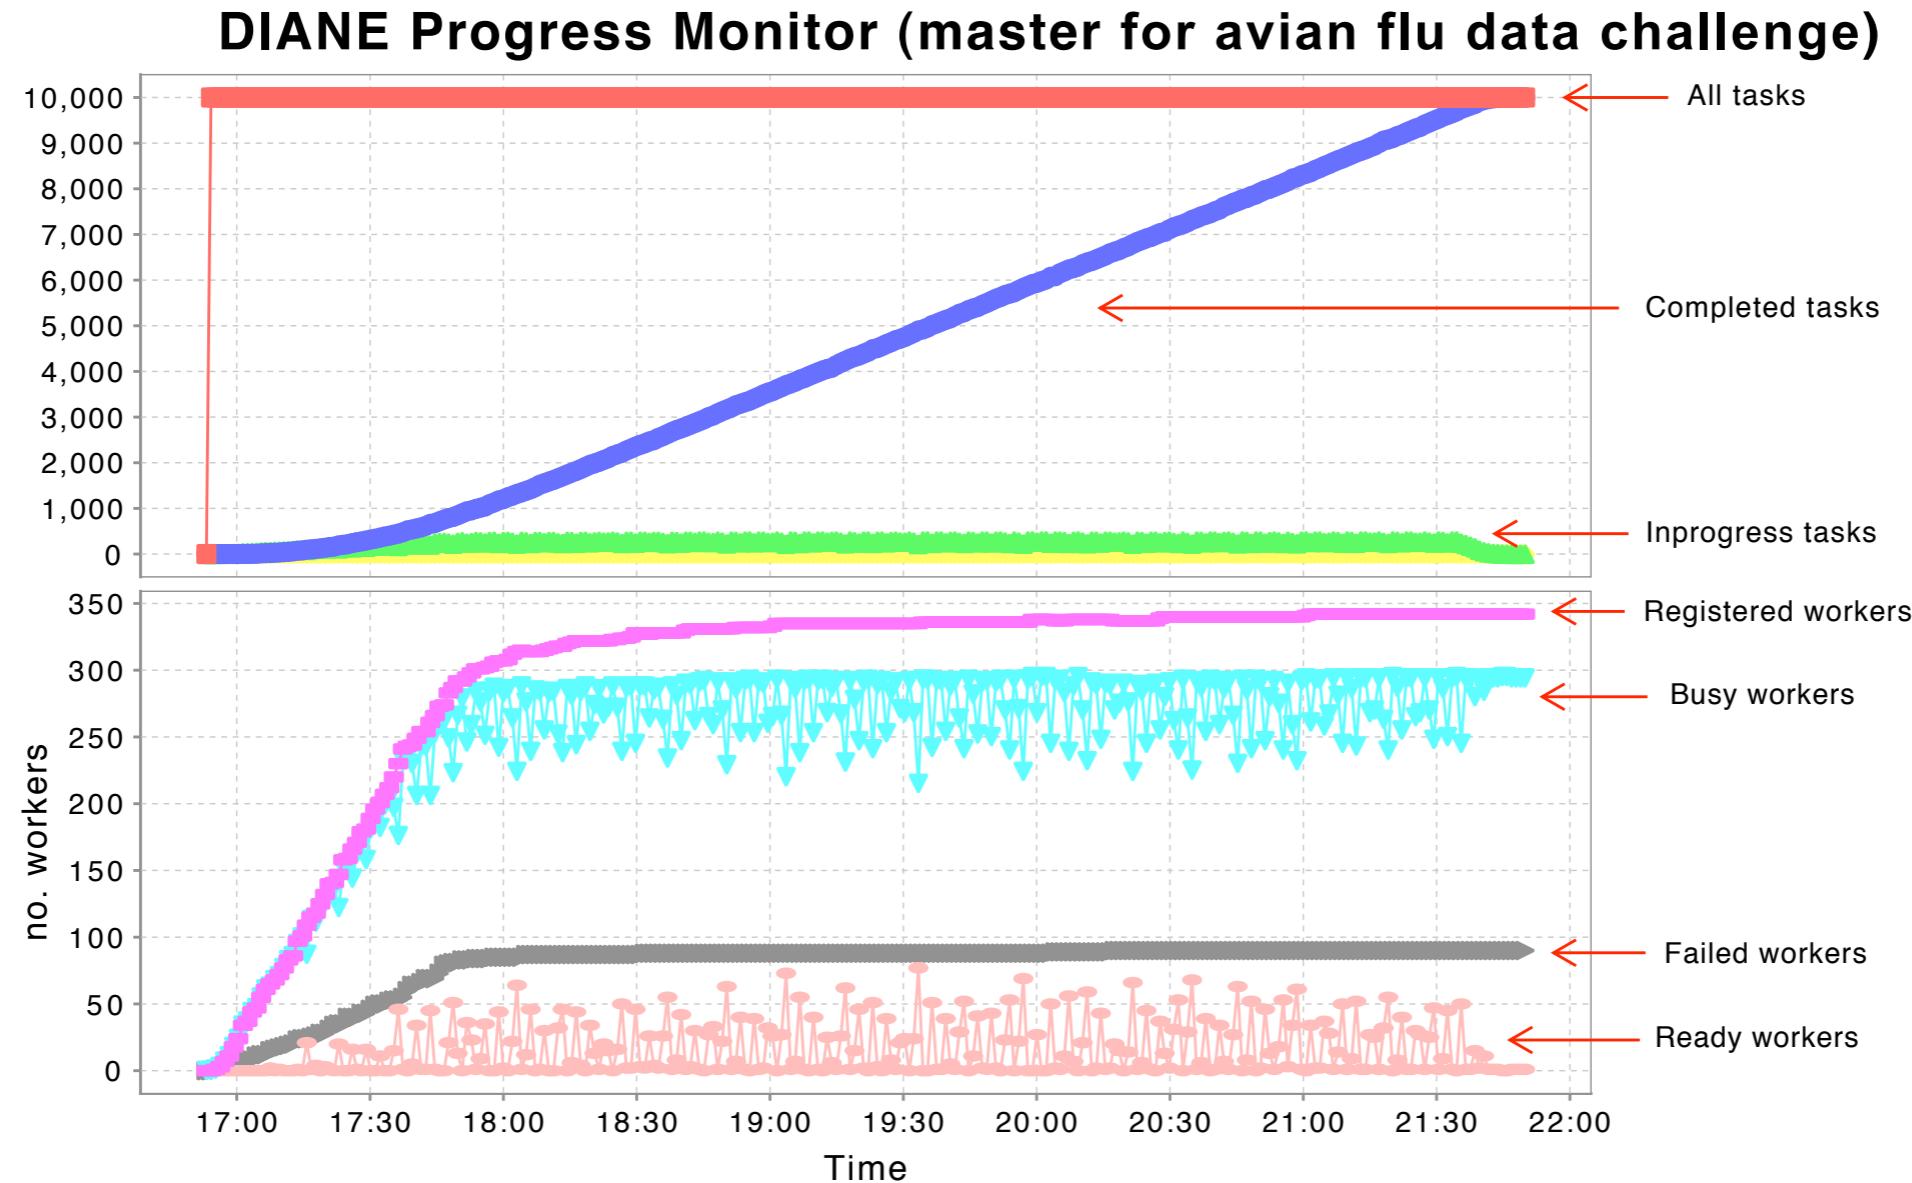


DIANE: drug discovery use case



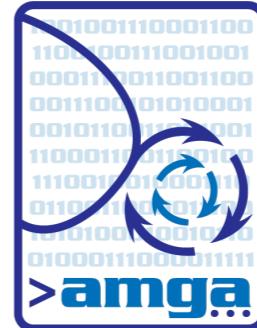
- > 100 distributed DIANE instances
- > 5000 CPUs
- > 1.2 Million simulations
- > 22 CPU years on 1 CPU
- < 1 months

DIANE: application reliability





AMGA: motivation



<http://cern.ch/ganga>

- Distributed data needs to be indexed
 - not only logical name \Leftrightarrow physical name (URI) mapping
 - but also high-level metadata (data about data)
- Relational database is another data storage
 - interoperability between different RDB technologies
 - intensive access through internet
- Data security, persistency and availability



AMGA: metadata concept

- Metadata - List of attributes associated with entries
 - Attribute – key/value pair with type information
 - Type – The type (int, float, string,...)
 - Name/Key – The name of the attribute
 - Value - Value of an entry's attribute
- Schema – A set of attributes
- Collection – A set of entries associated with a schema
- Think of schemas as tables, attributes as columns, entries as rows



AMGA: metadata catalogue features

- Dynamic Schemas
 - Schemas can be modified at runtime by client
 - Create, delete schemas
 - Add, remove attributes
- Metadata organized as an hierarchy
 - Collections can contain sub-collections
 - Analogy to file system:
 - Collection ⇔ Directory
 - Entry ⇔ File
- Flexible Queries
 - SQL-like query language
 - Joins between schemas

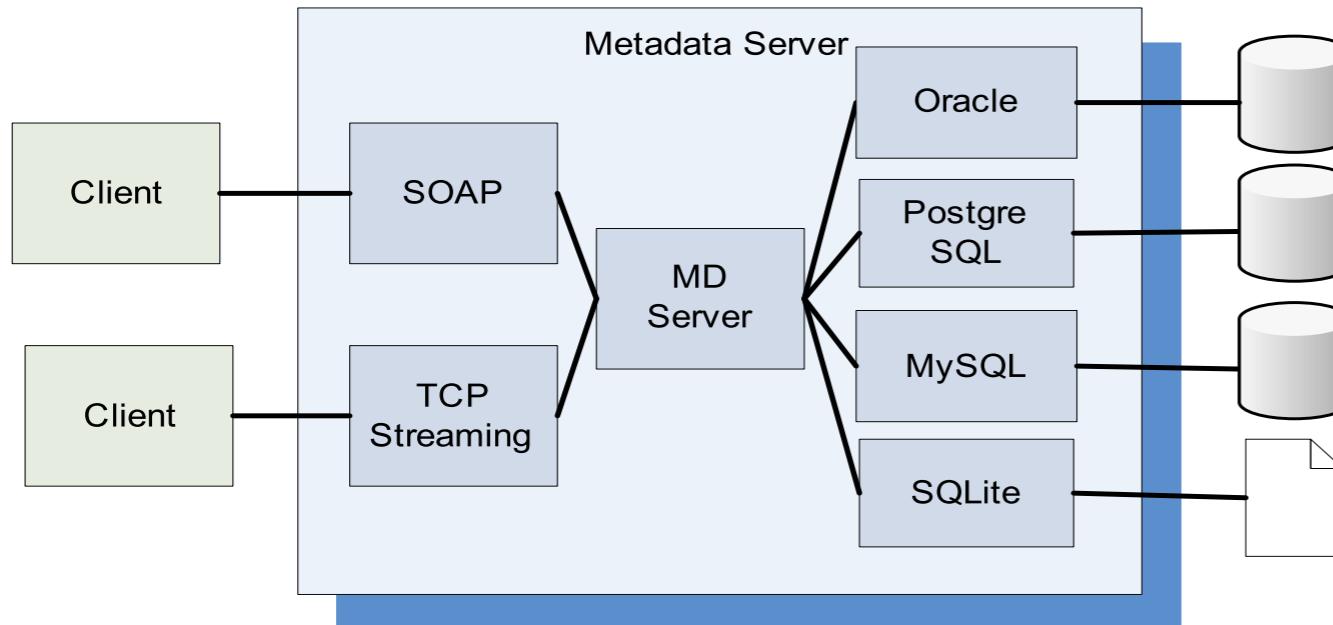
Supporting filesystem-like operations:

ls
mv
mkdir

QUERY EXAMPLE:

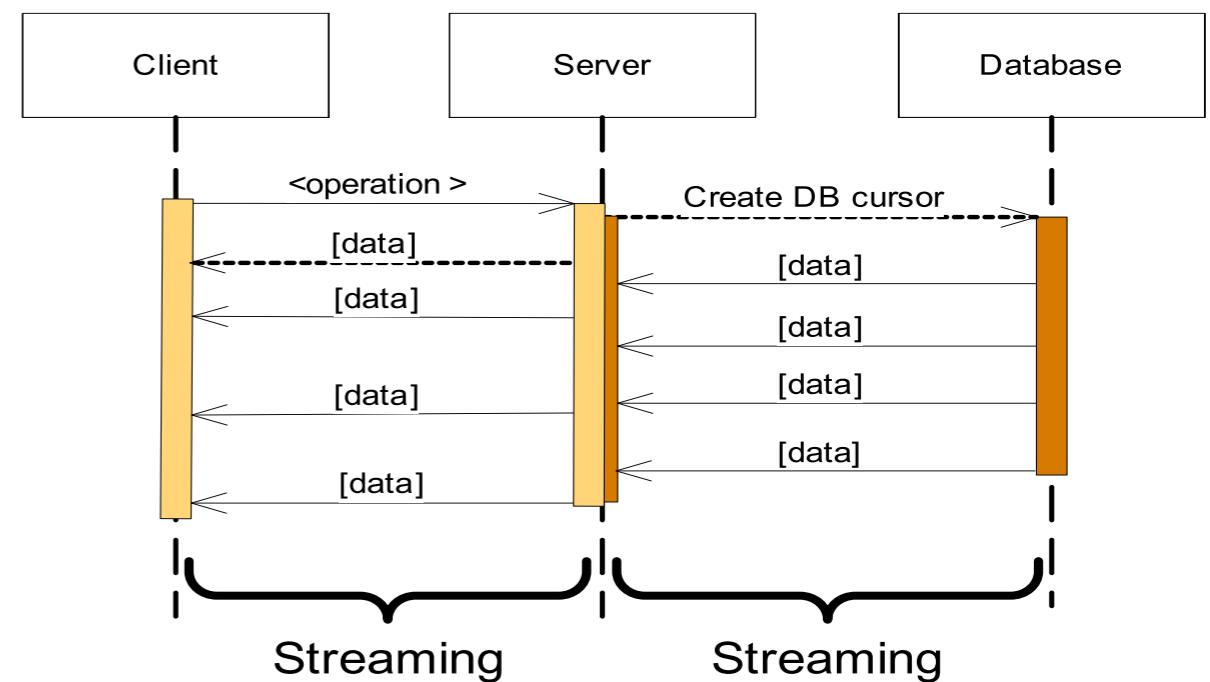
```
selectattr /gLibrary:FileName \
/gLibrary:Author \
'/gLibrary:FILE=/gLAudio:FILE \
and \
like(/gLibrary:FileName, "% .mp3")'
```

AMGA: client-server model



- RDB neutral
- TCP streaming for performance
- SOAP for adoptability

- Asynchronous streaming reducing operation overhead



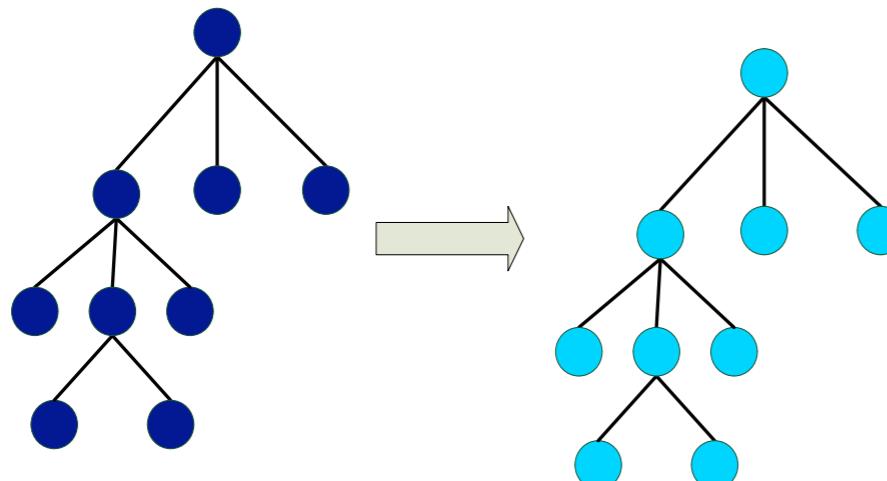


AMGA: security

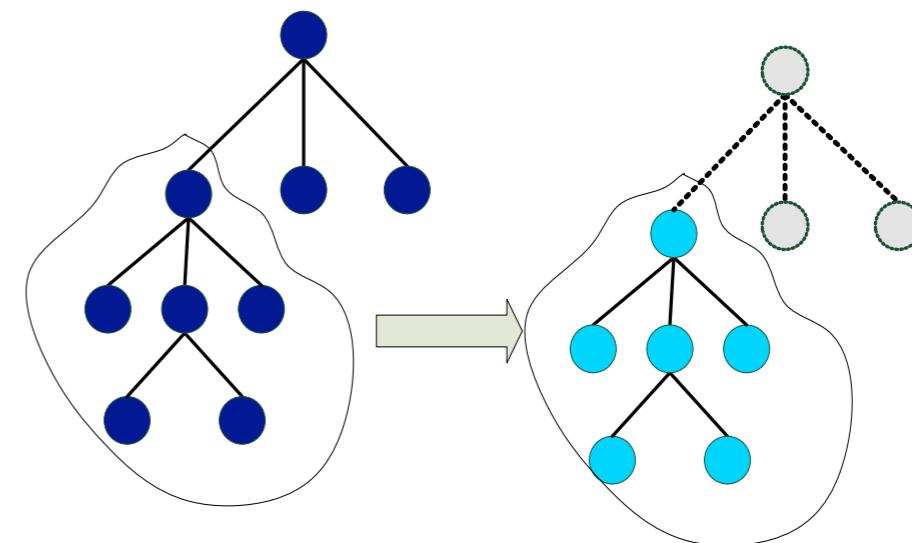
- Unix style permissions
- ACLs – per-collection or per-entry
- Secure connections – SSL
- Client Authentication based on
 - Username/password
 - General X509 certificates
 - Grid-proxy certificates
- Access control via a Virtual Organization Management System (VOMS)

AMGA: metadata preservation / replication

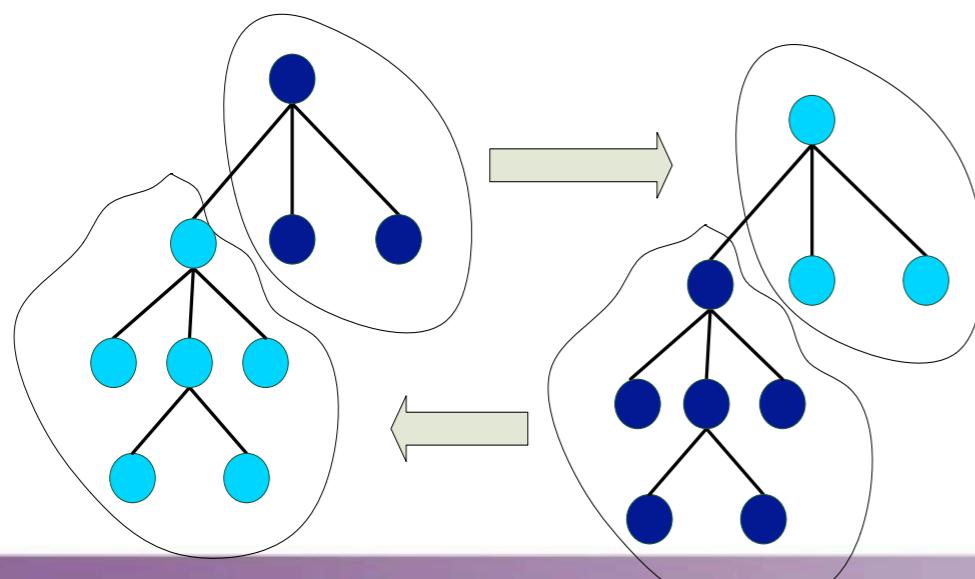
Full replication



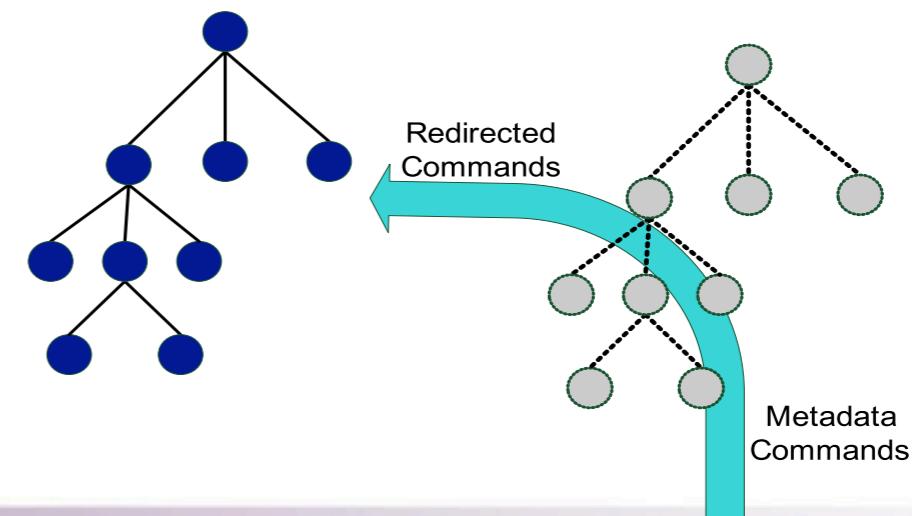
Partial replication



Federation



Proxy





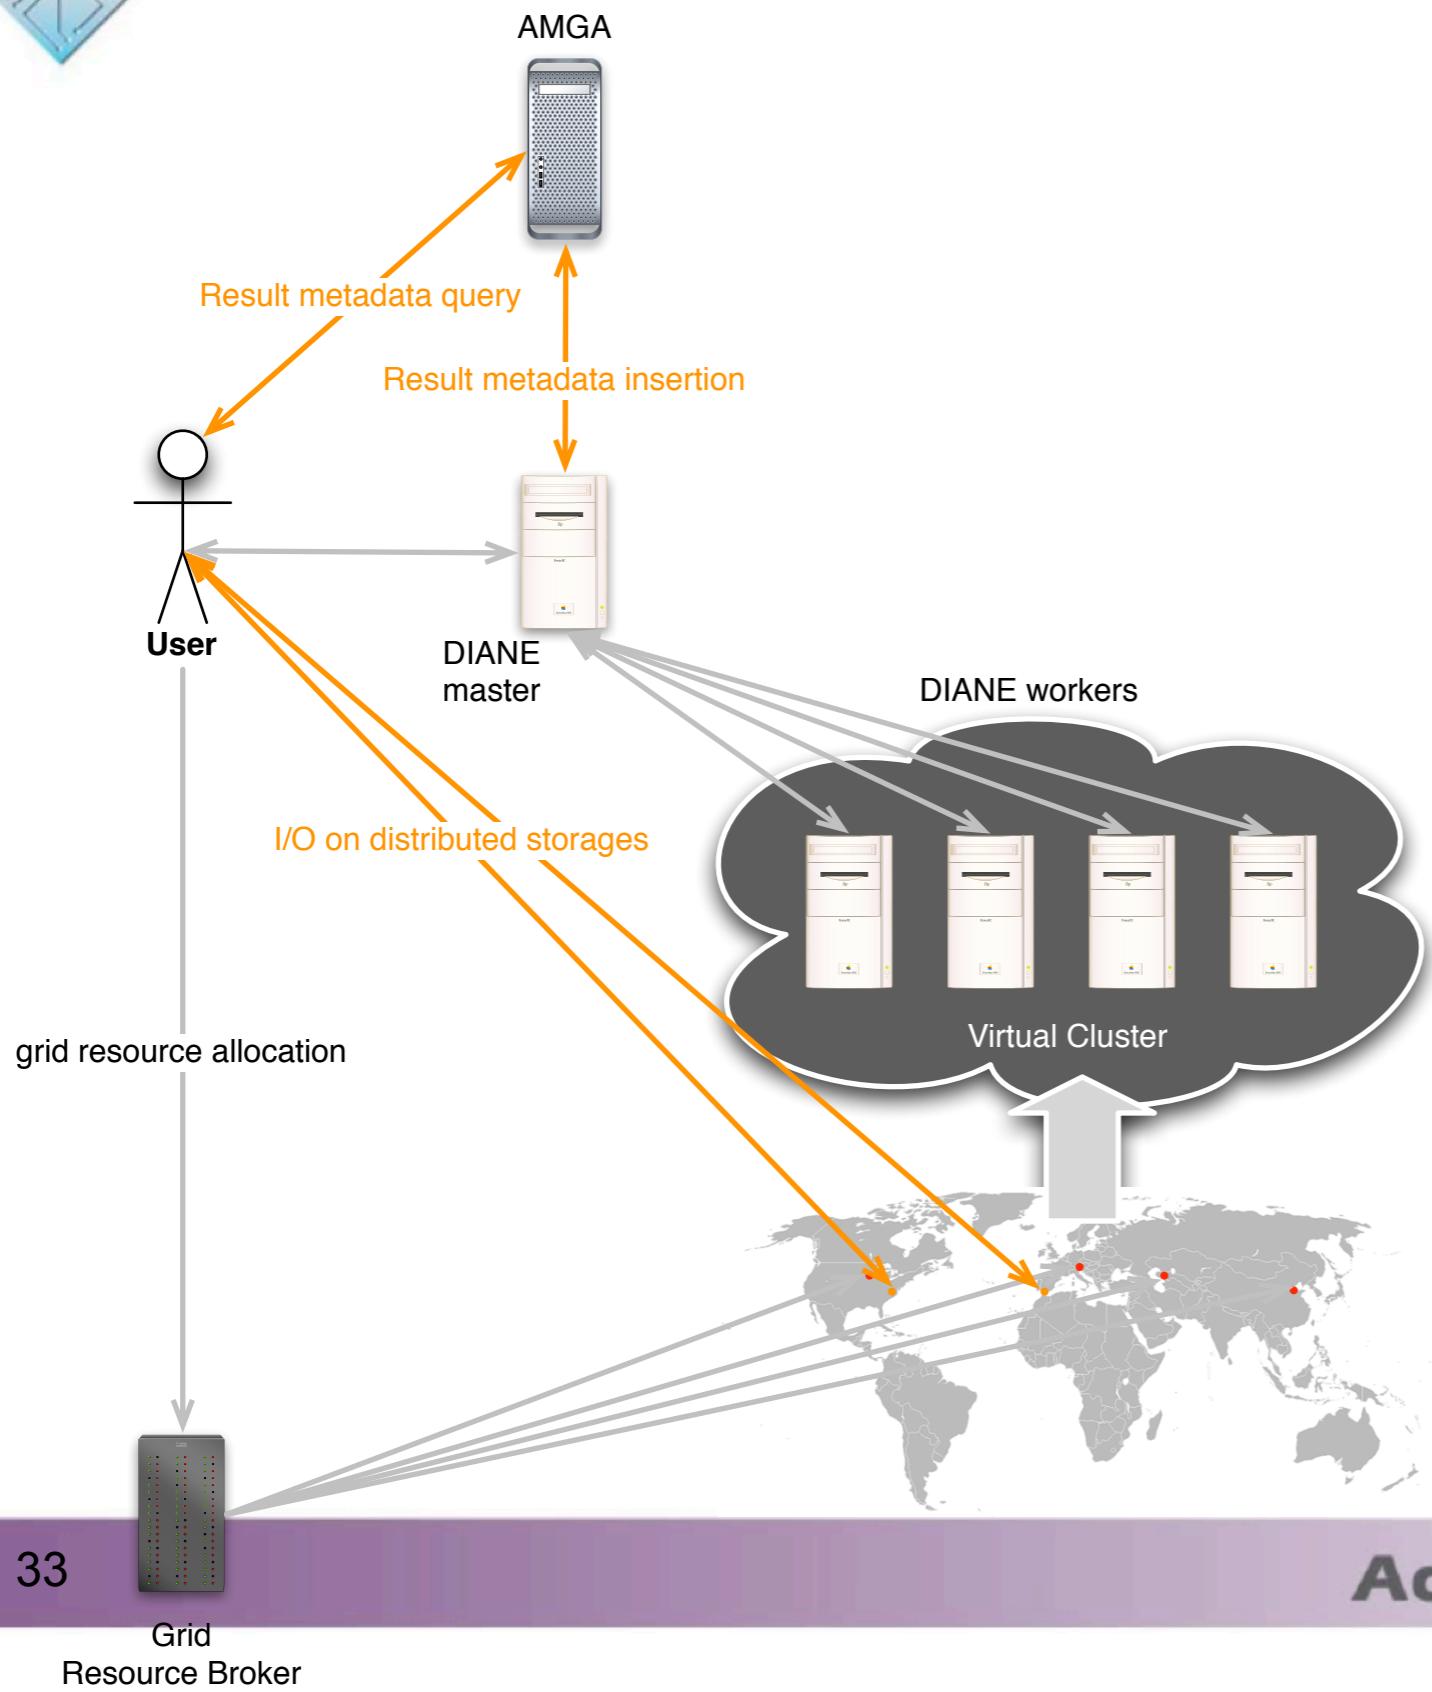
AMGA: drug analysis use case

- Questions:
 - large-scale analysis is a long process; can I investigate my result of the completed simulations during the runtime?
 - there are too many simulation results distributed on several storages; can I just access some of them that I have interests (according to some criteria)?
- Answer:

YES! We need to indexing results with their essential attributes ... ⇒ Metadata catalogue

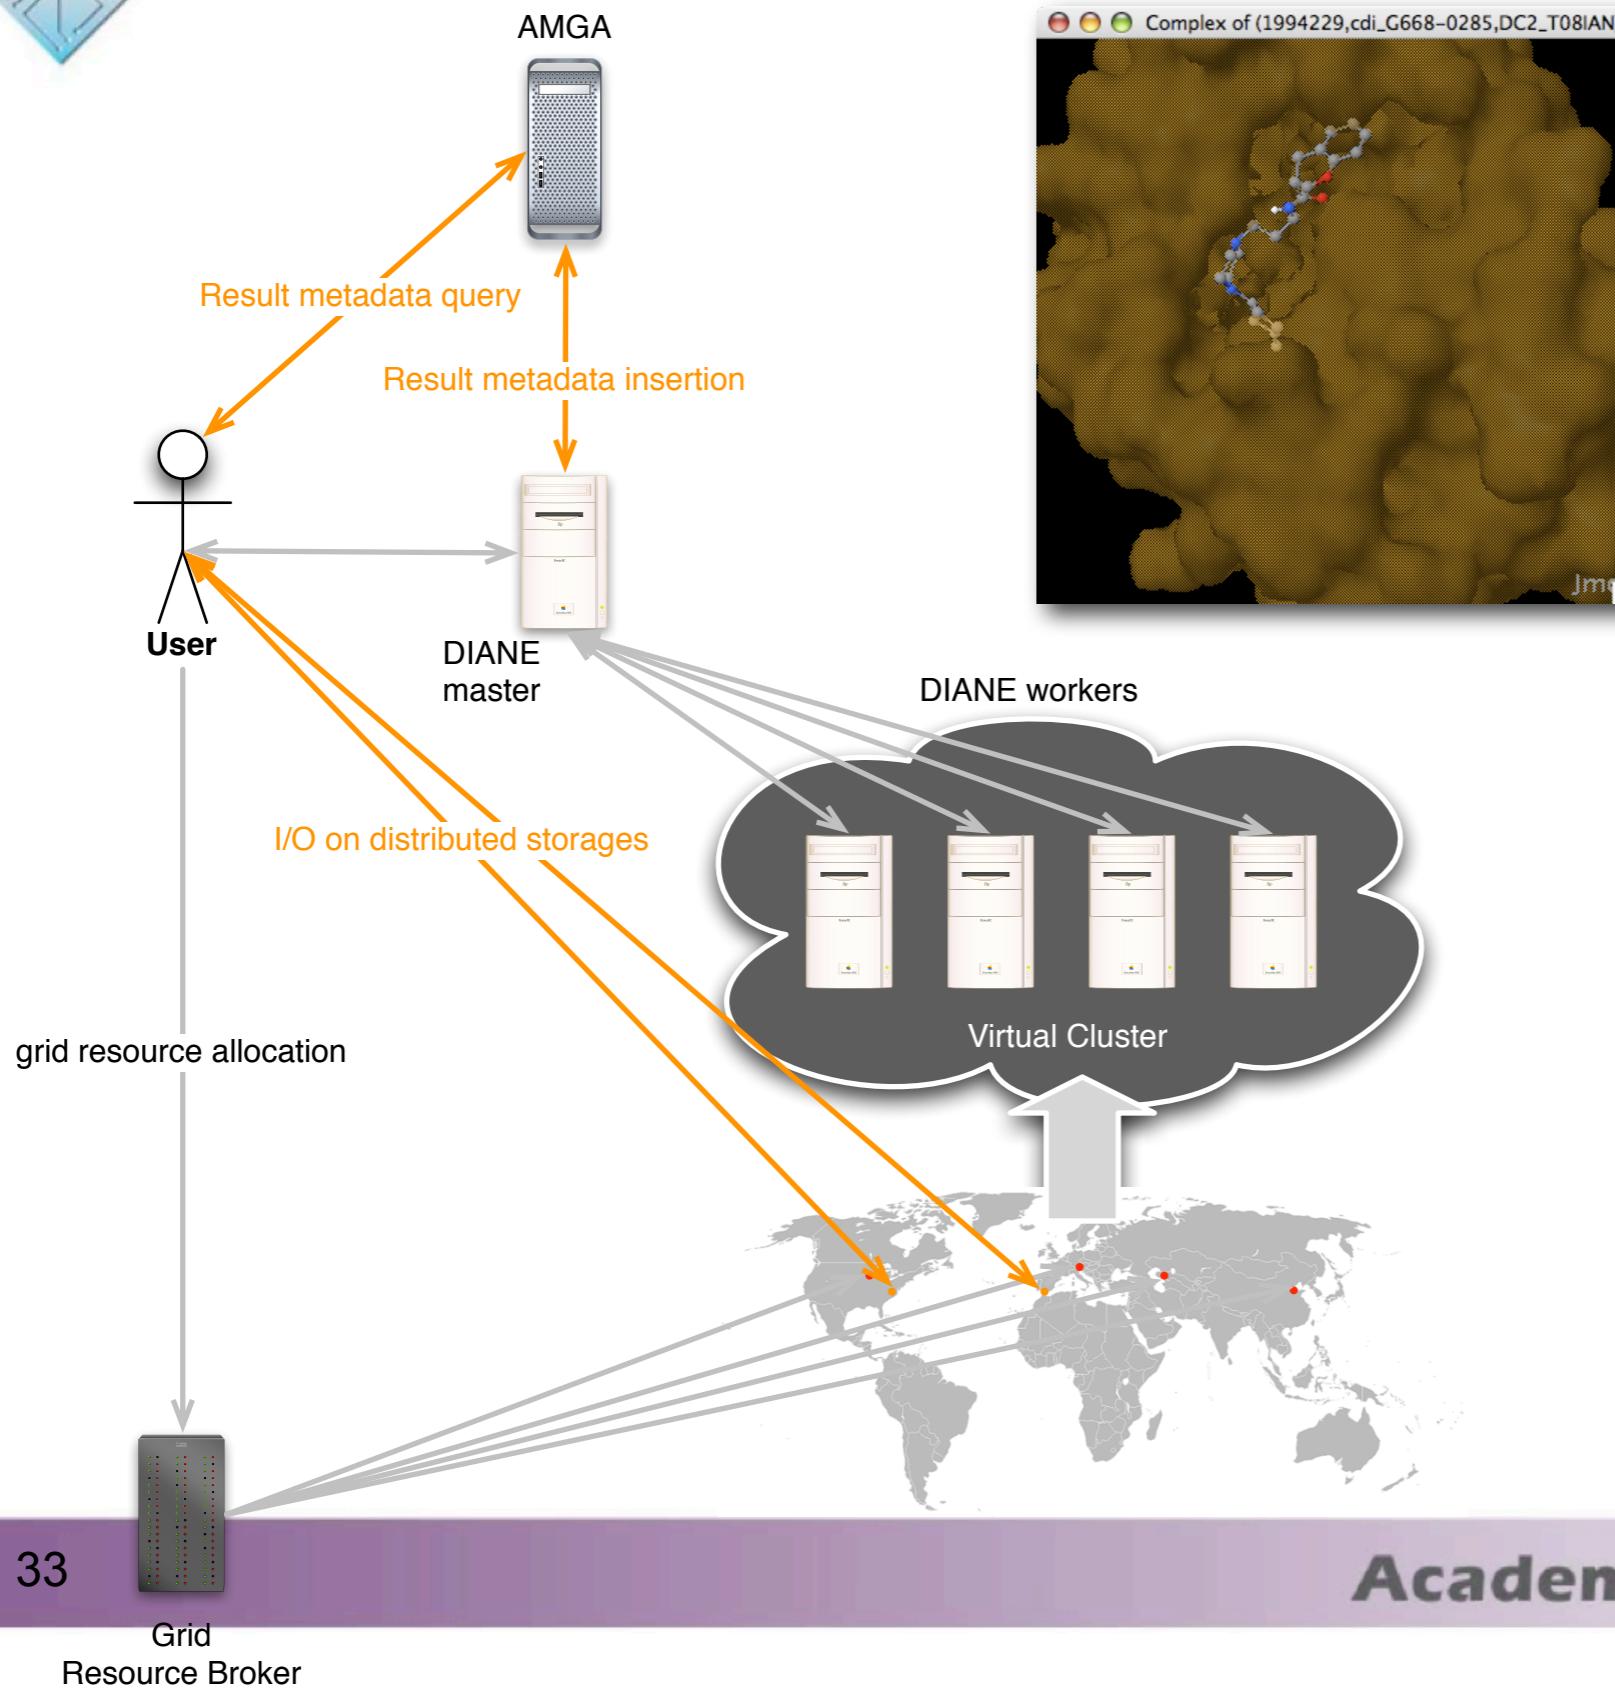


AMGA: drug analysis use case



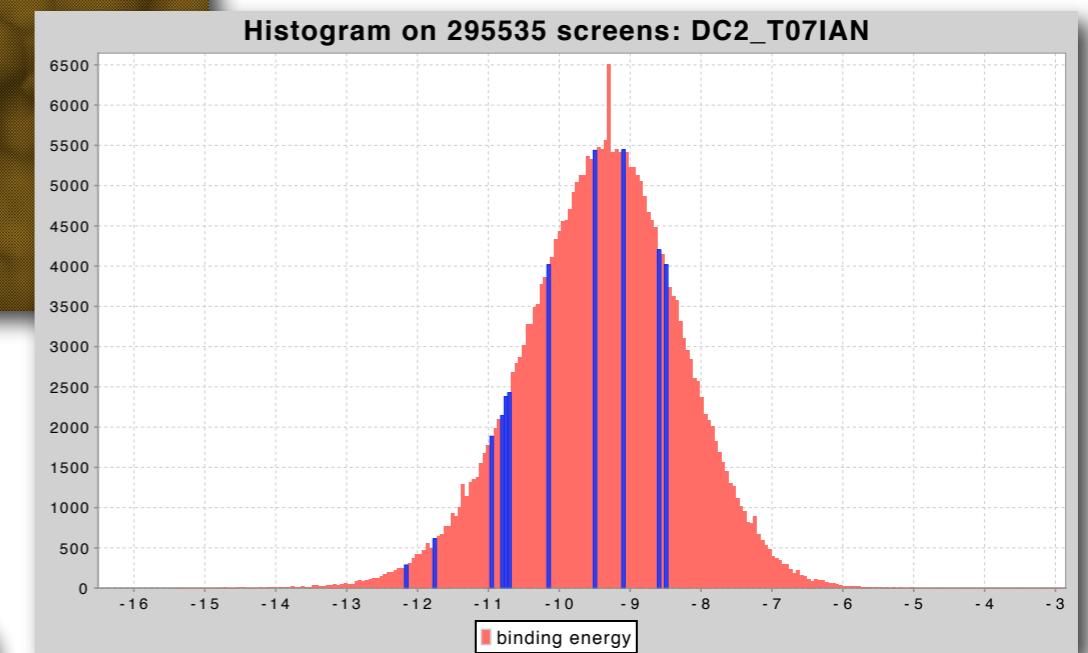
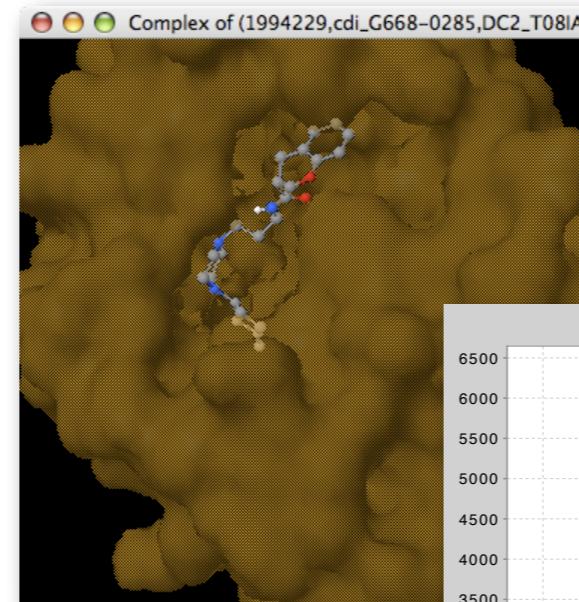
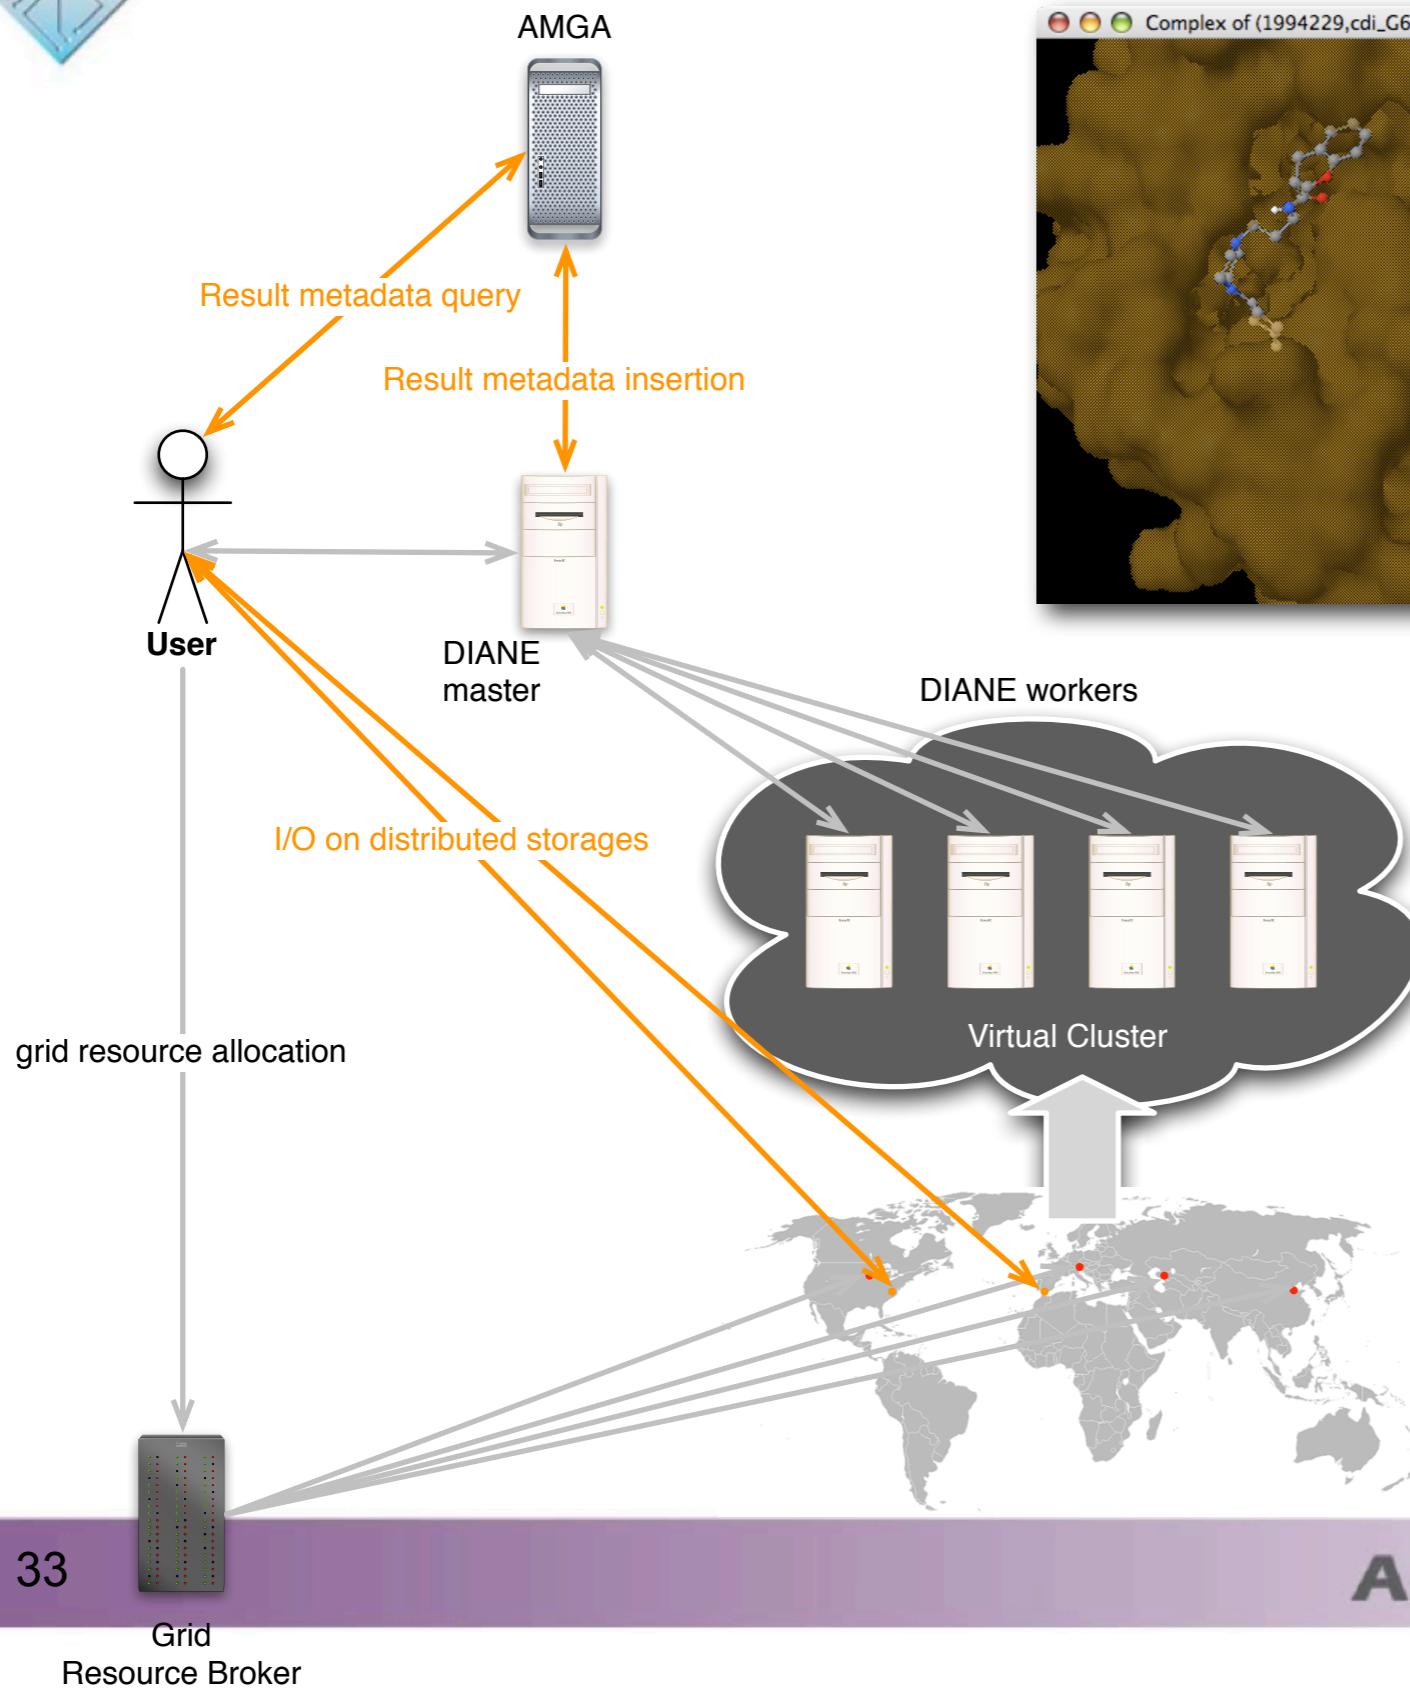


AMGA: drug analysis use case



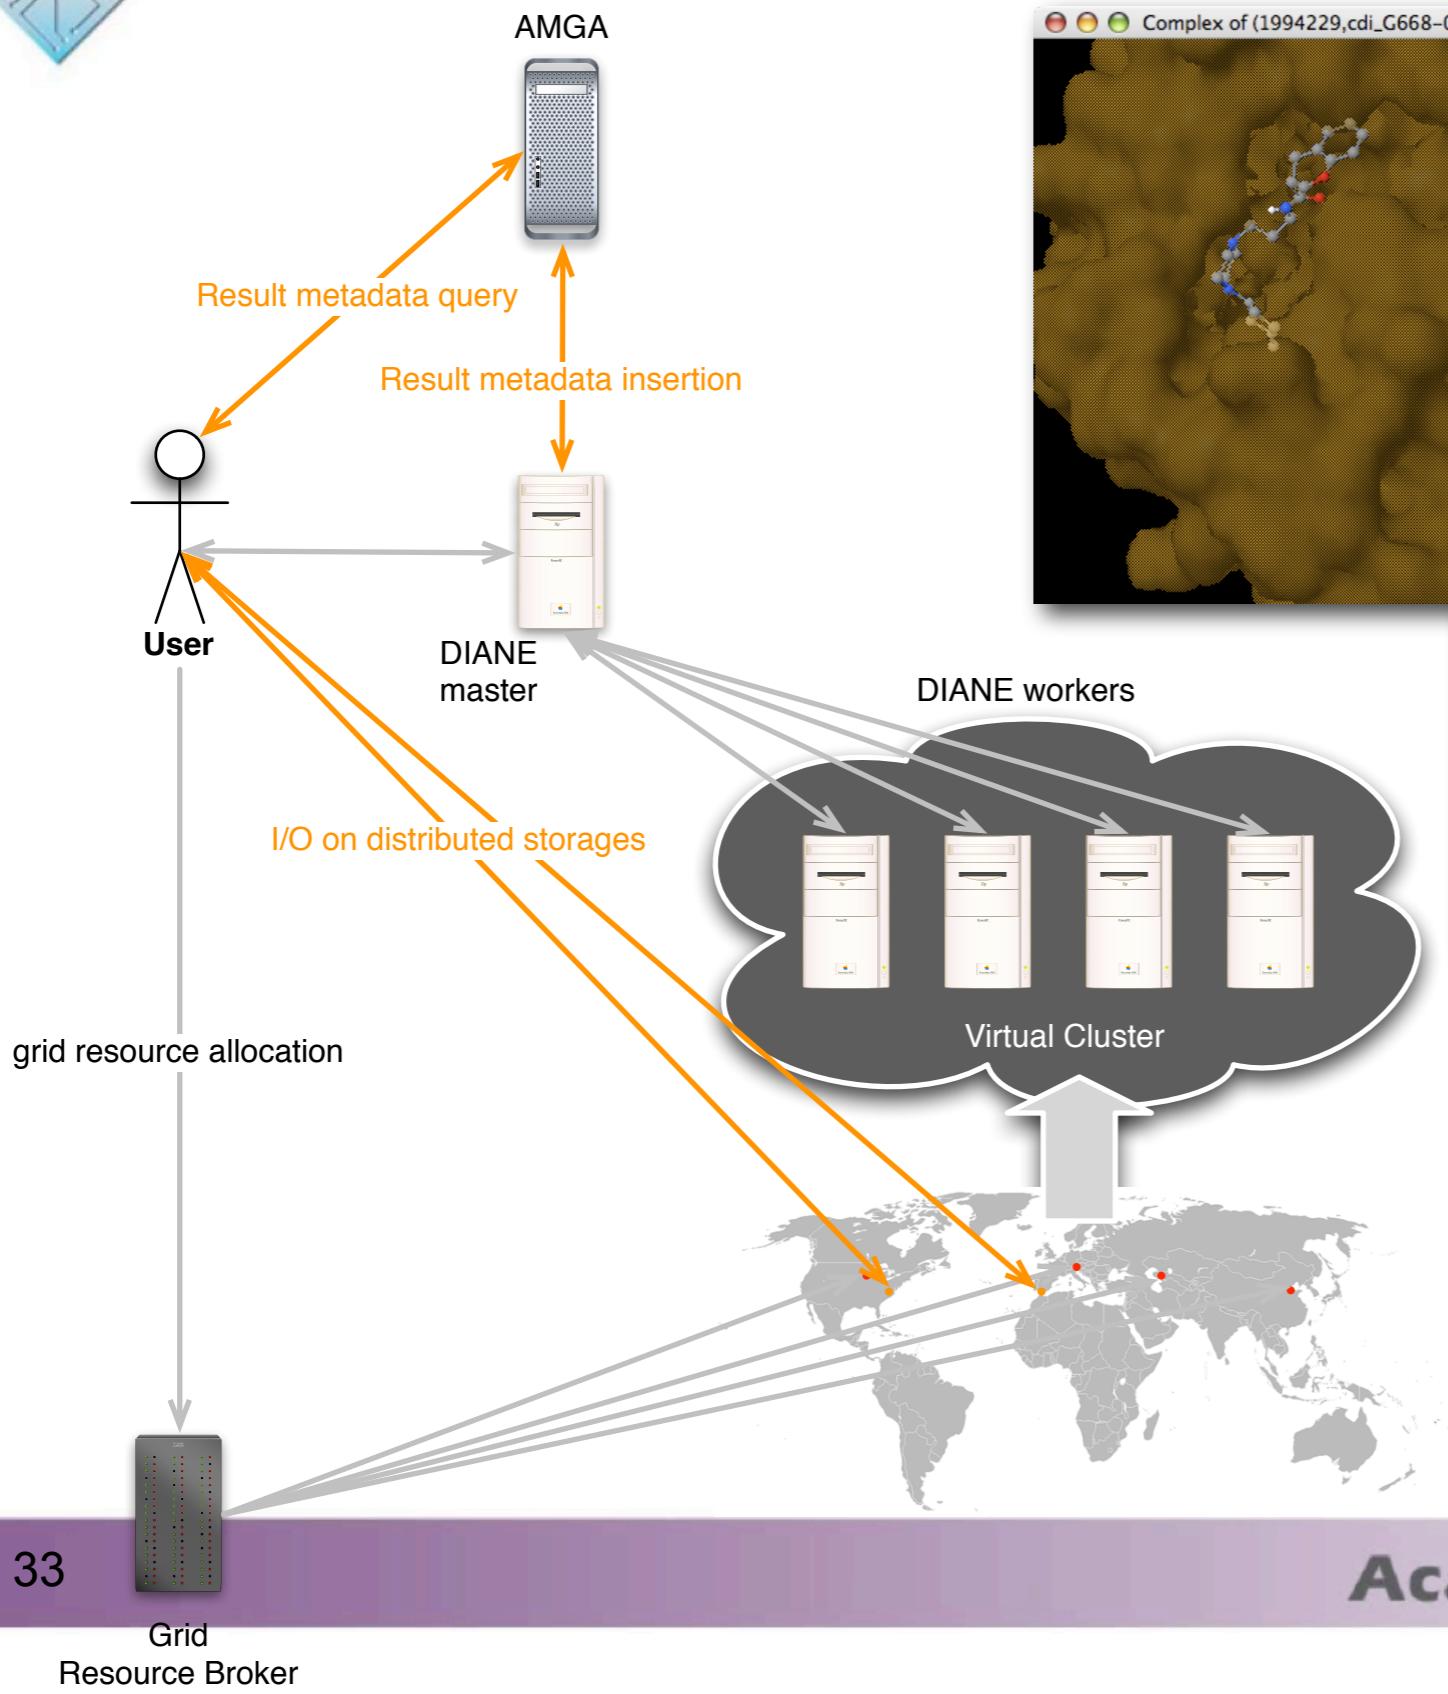


AMGA: drug analysis use case





AMGA: drug analysis use case





Conclusion

- Grid is about computing resource sharing
- Grid development is a software engineering
 - working on top of the existing computing infrastructure
 - driven by applications
- High-level utilities are needed in order to fill up the gap between the applications and the middleware
 - where the open source community could contribute



Thank you!

TWGRID: www.twgrid.org